

Topic 07:

Multi-scale image transforms

- Gaussian & Laplacian pyramids
- Linear image transforms
- Haar wavelet transform

Topic 07:

Multi-scale image transforms

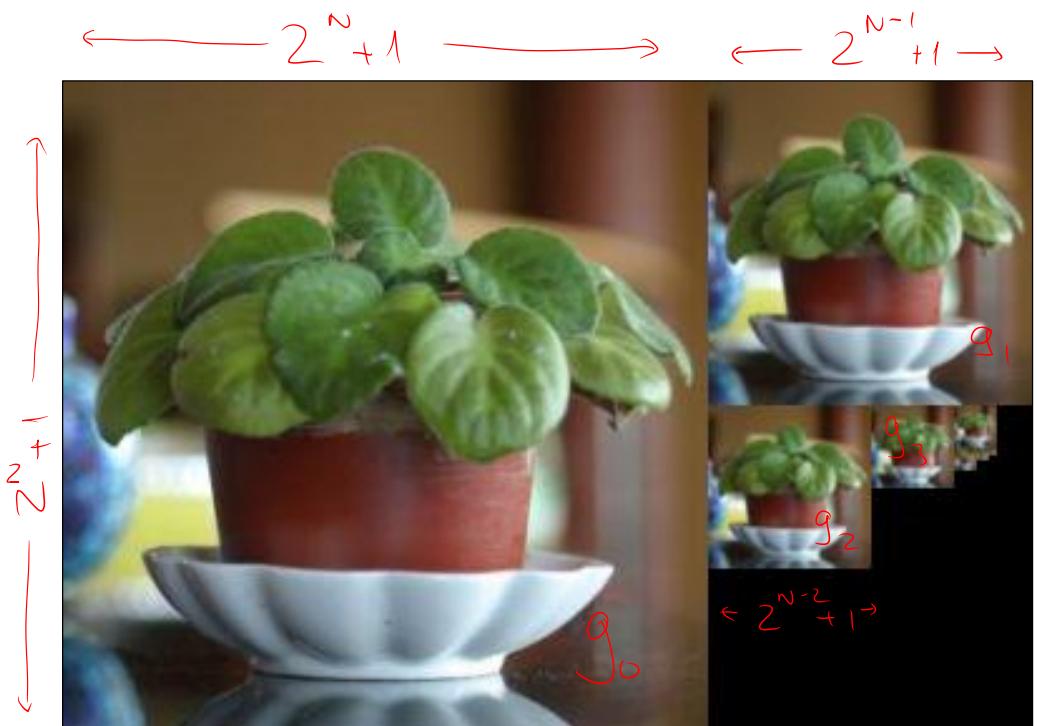
- Gaussian & Laplacian pyramids
- Linear image transforms
- Haar wavelet transform

Gaussian pyramid

Goal: Develop representation to decompose images into information at multiple scales, to extract features or structures of interest, to attenuate noise

Input:

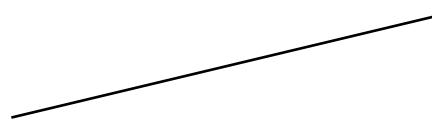
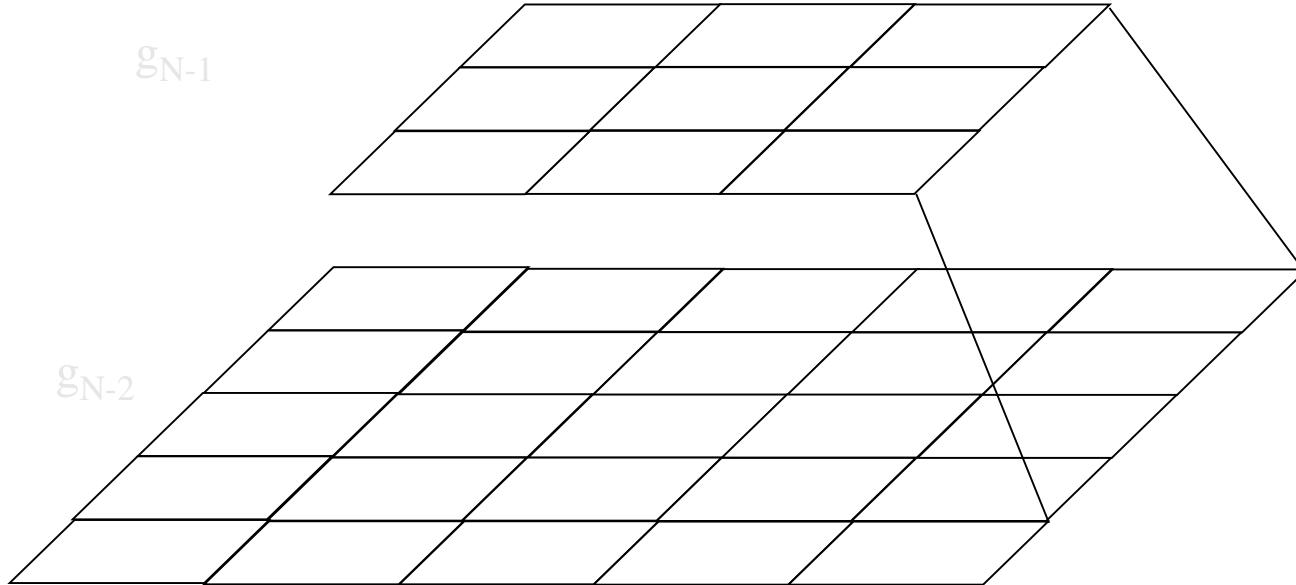
Image I of size
 $(2^N+1) \times (2^N+1)$



Output:

- N images g_0, \dots, g_{N-1}
- g_e has size
 $(2^{N-l}+1) \times (2^{N-l}+1)$

why a “pyramid”?



g_0 (= original image)



Gaussian pyramid

The representation is based on 2 basic operations

① Smoothing

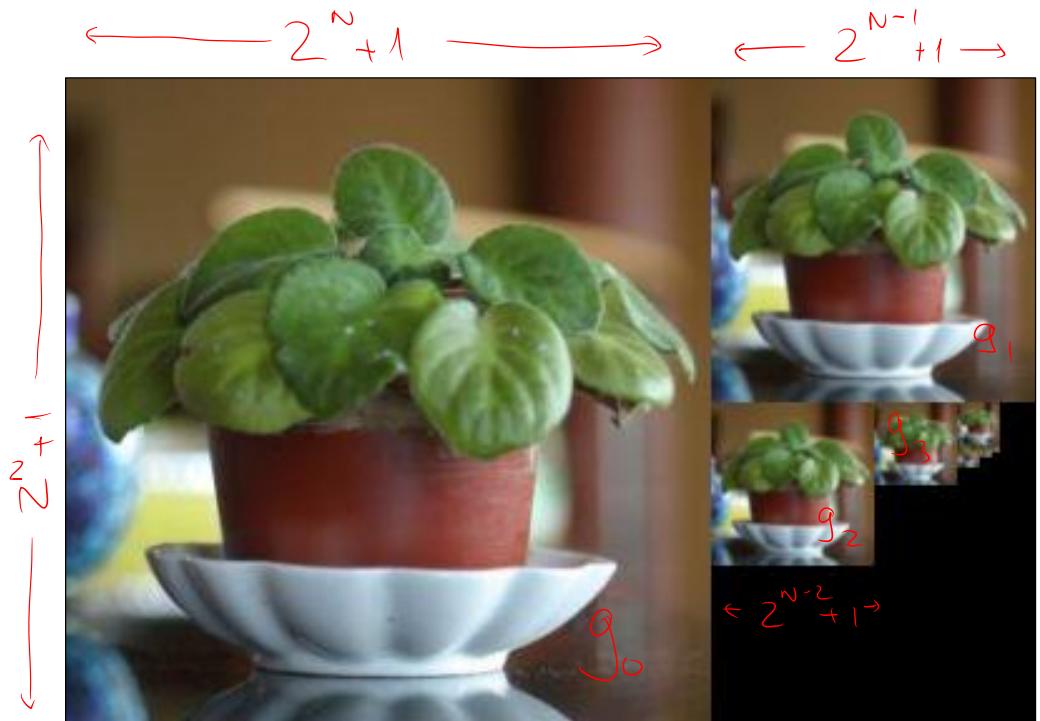
Smooth the image with a sequence of smoothing filters each of which has twice the radius of the previous one

② Downsampling

Reduce image size by $1/2$ after each smoothing

Aside: Downsampling = any linear transformation

$$\xrightarrow{\text{downsampled image}} \left[\begin{array}{c} \text{matrix} \\ \# \text{rows} \\ \# \text{columns} \end{array} \right] \left[\begin{array}{c} \text{original image} \end{array} \right]$$



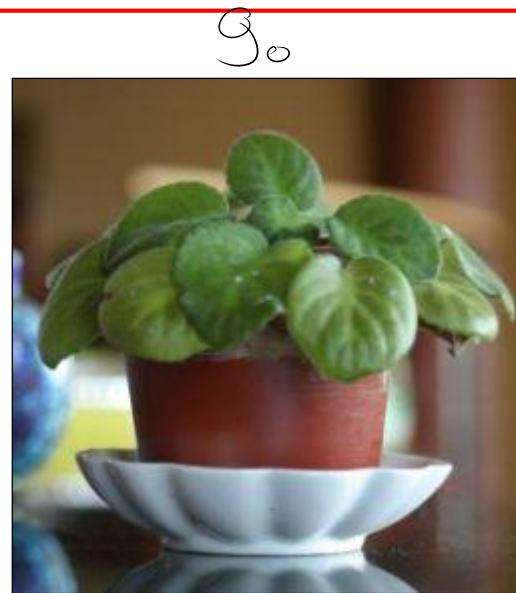
operation #1: recursive smoothing with filter w

Original photo g_0

$$\hat{g}_1 = w * g_0$$



a 5×5
filter



$$\hat{g}_1(i, j) = \sum_{m=-2}^2 \sum_{n=-2}^2 w(m, n) \cdot g_0(i-m, j-n)$$

operation #1: recursive smoothing with filter w

$$\hat{g}_2 = w * \hat{g}_1$$

$$= w * (w * g_0)$$

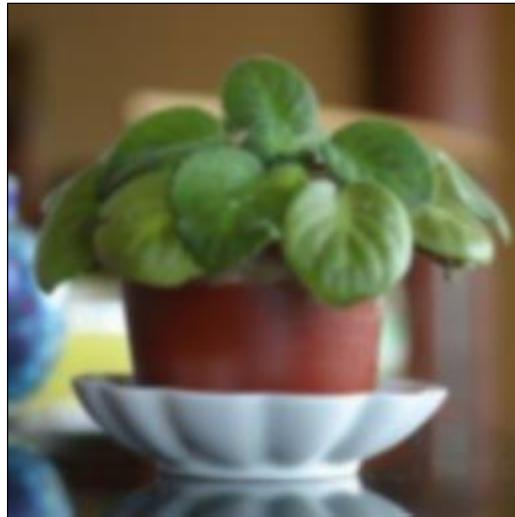
$$= \underbrace{(w * w)}_{\text{can be thought of as a filter}} * g_0$$

can be thought of
as a filter

$$h = w * w$$

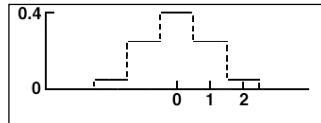
whose radius is
twice that of w

\hat{g}_1

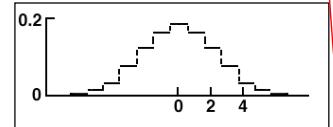


In 1D:

w:



w * w:



operation #1: recursive smoothing with filter w

$$\hat{g}_3 = w * \hat{g}_2$$

$$= \underbrace{(w * w * w)}_{\text{radius is 4 times that of } w} * g_0$$

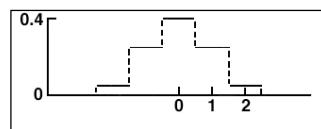
radius is 4
times that of w

\hat{g}_2

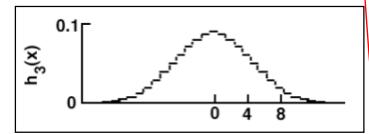


In 1D:

w :



$w * w * w$:

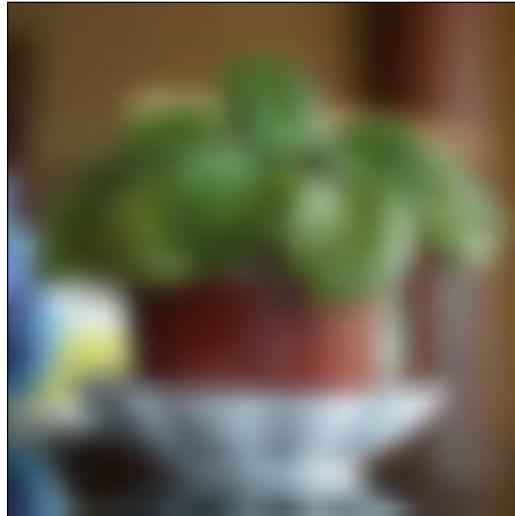


operation #1: recursive smoothing with filter w

$$\hat{g}_4 = w * \hat{g}_3$$

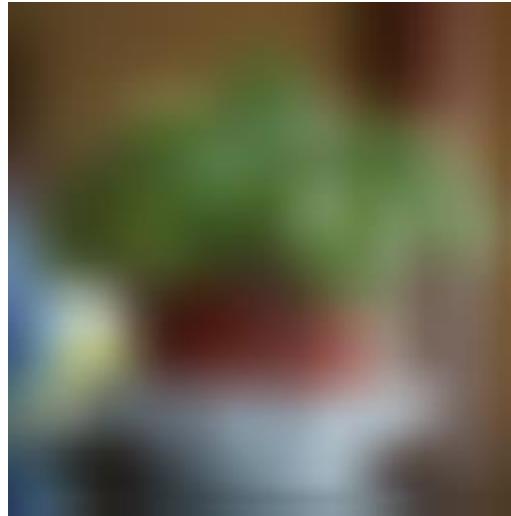
$$= (w * w * w * w) * g_0$$

\hat{g}_3



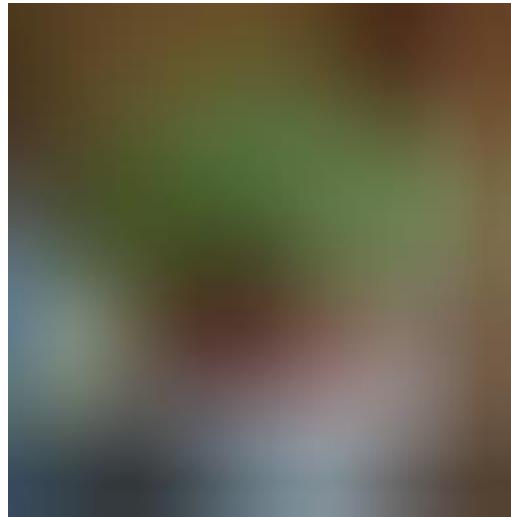
operation #1: recursive smoothing with filter w

\hat{g}_4



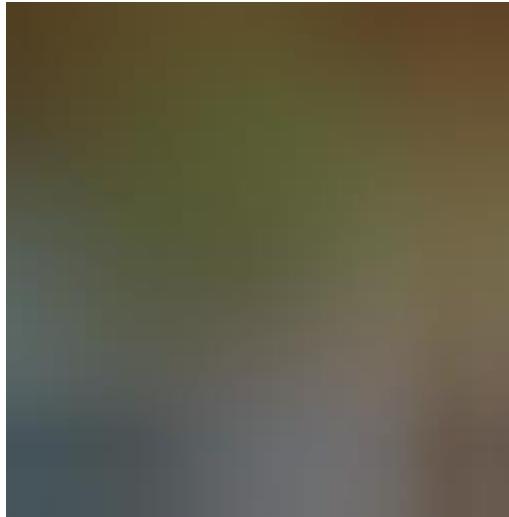
operation #1: recursive smoothing with filter w

\hat{g}_s



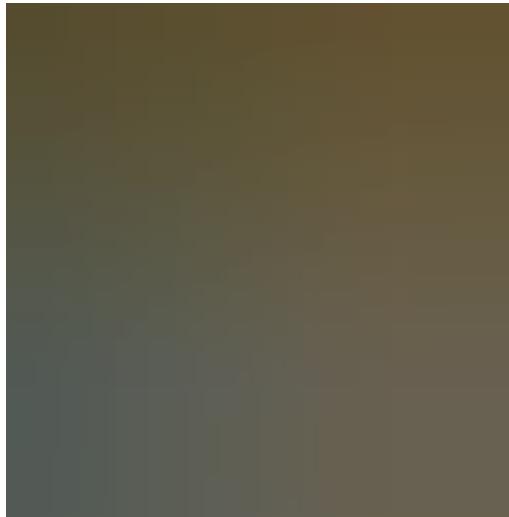
operation #1: recursive smoothing with filter w

$$\hat{\mathcal{G}}_6$$



operation #1: recursive smoothing with filter w

\hat{g}_7



deriving the smoothing filter in 1D

① \hat{w} always has s elements
 (a.k.a "s-tap" filter)

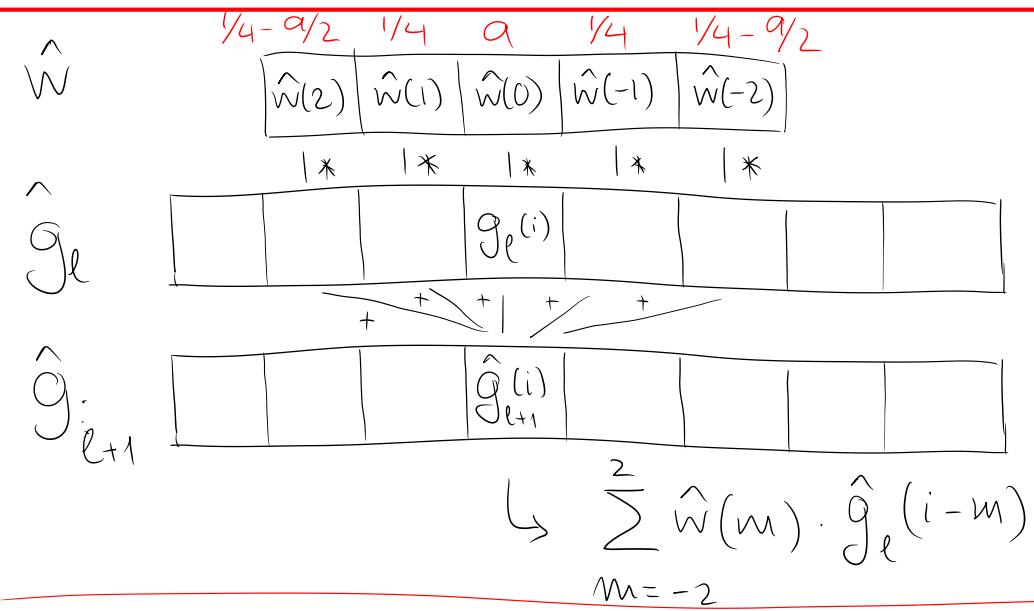
② \hat{w} symmetric about 0:
 $\hat{w} = \begin{bmatrix} c \\ b \\ a \\ b \\ c \end{bmatrix} = \begin{bmatrix} \frac{1}{4}-\frac{a}{2} \\ \frac{1}{4} \\ a \\ \frac{1}{4} \\ \frac{1}{4}-\frac{a}{2} \end{bmatrix}$

③ applying \hat{w} to a constant image does not change it

$$\Leftrightarrow \sum_{m=-2}^2 \hat{w}(m) = 1$$

$$m=-2$$

$$\Leftrightarrow a + 2b + 2c = 1$$



④ Equal contribution.

$$a + 2b = 2b = \frac{1}{2}$$

To satisfy Criteria 1-4 we have 2 eqs & 3 unknowns
 $\Rightarrow a$ remains free parameter

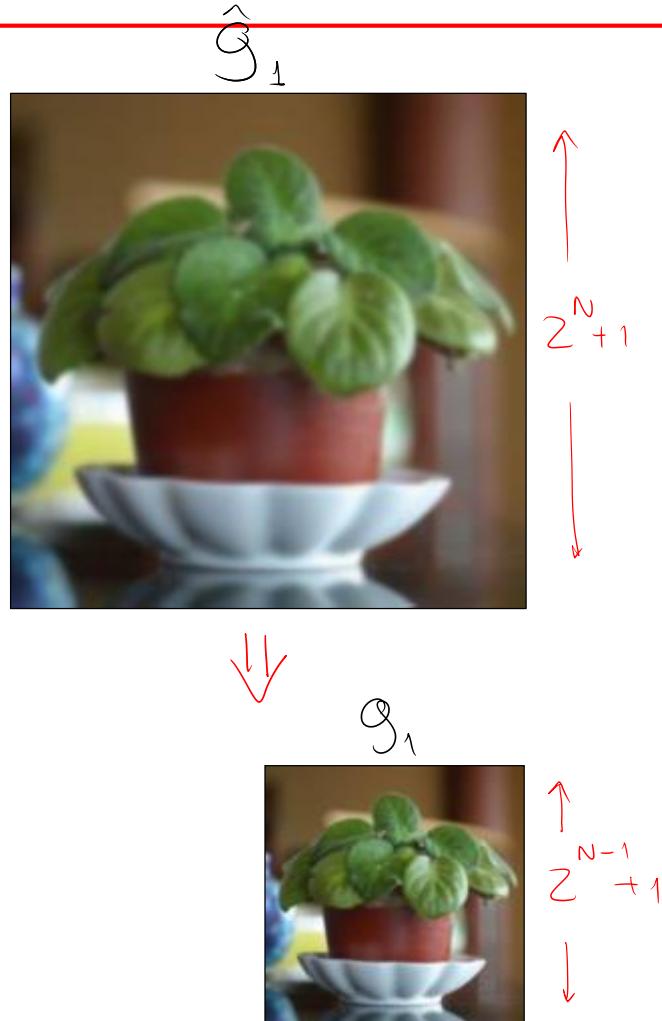
$$\hat{w}(2) = \hat{w}(-2) = \frac{1}{4} - \frac{a}{2}, \hat{w}(-1) = \hat{w}(1) = \frac{1}{4}$$

$$\text{usually } a \in [0.3, 0.6]$$

operation #2: recursive downsampling

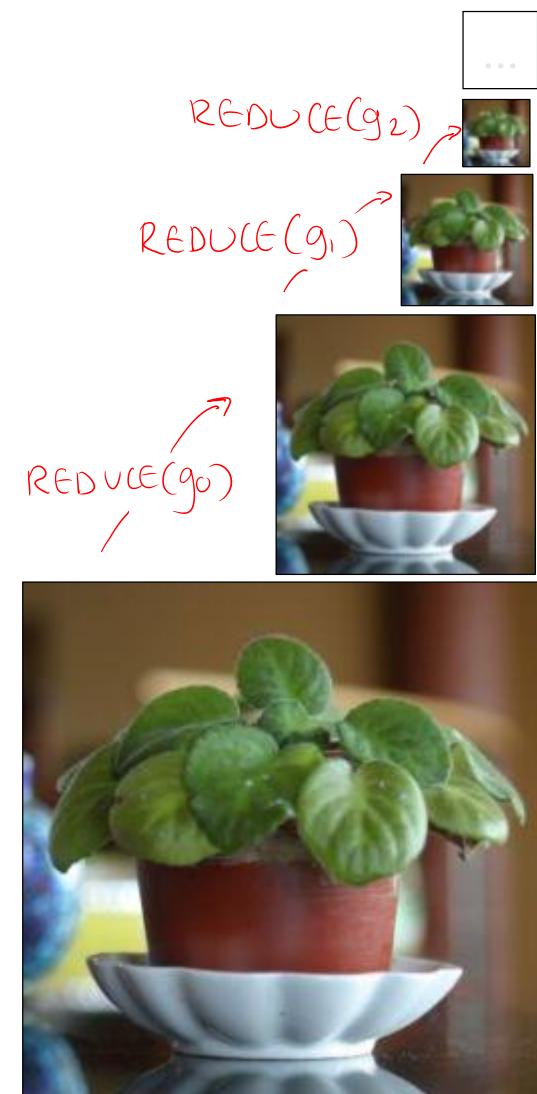
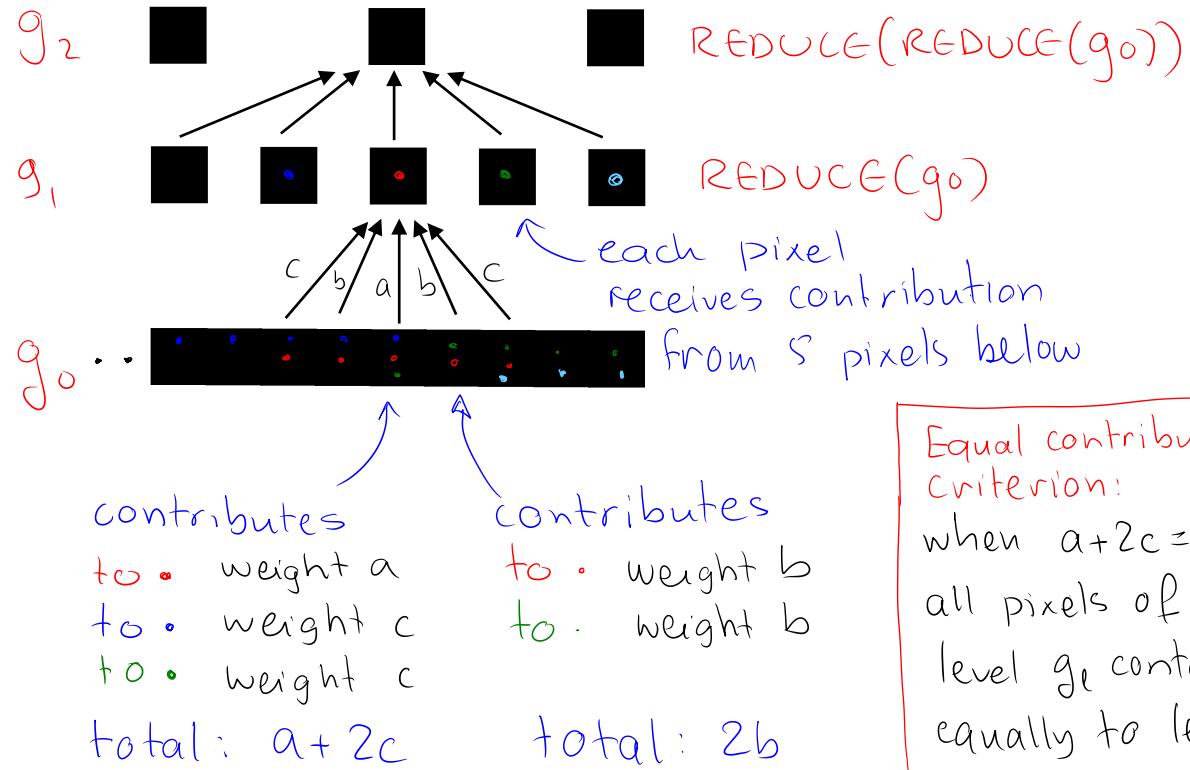
Since \hat{g}_1 contains less image detail, we downsample it by 2 (i.e. store every other pixel)

$$g_1(i,j) = \hat{g}_1(2i, 2j)$$



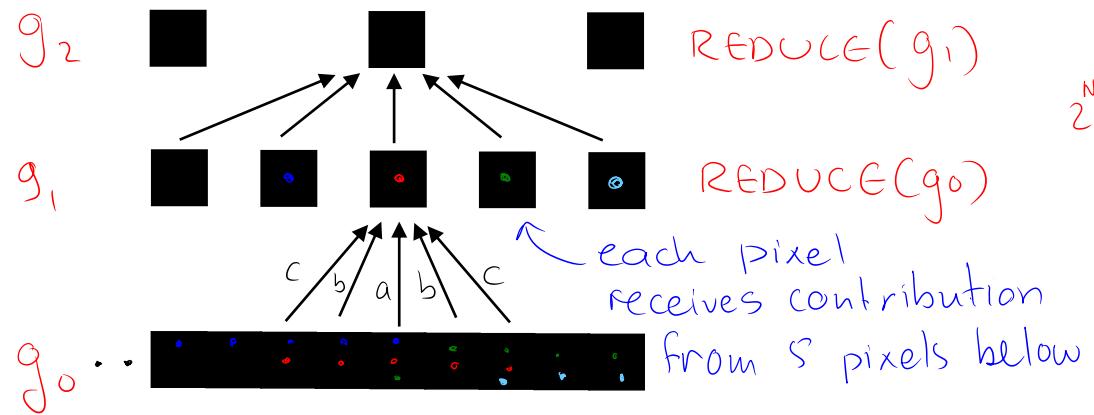
the REDUCE() function

Visualizing in 1D



the REDUCE() function

Visualizing in 1D



REDUCE function in matrix notation (1D)

$$g_1 = D_o \cdot C_o \cdot g_0$$

downsampling matrix D_o

convolution matrix C_o

$\left[\begin{array}{c} g_1 \\ \vdots \end{array} \right] = \left[\begin{array}{ccccccccc} 1 & 0 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ 0 & 0 & 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \ddots \end{array} \right] \left[\begin{array}{ccccccccc} c & b & a & b & c & \dots & c & b & a \\ b & a & b & c & \dots & b & a & b & c \\ a & b & c & \dots & a & b & c & \dots & a \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \ddots & \ddots & \vdots \\ 0 & \dots & 0 & \dots & 0 & \dots & 0 & \dots & 0 \end{array} \right] \left[\begin{array}{c} g_0 \\ \vdots \end{array} \right]$

$$g_1 = D_o \cdot C_o \cdot g_0$$

General expression:

$$g_{l+1} = D_l \cdot C_l \cdot g_l$$

smoothing in 2D

$$\hat{g}_1 = w * g_0$$

w is a separable smoothing filter defined by \hat{w}

$$\begin{array}{c} \uparrow s \\ \text{grid} \\ \downarrow s \\ \leftarrow s \quad \rightarrow \end{array} = \begin{bmatrix} \hat{w} \end{bmatrix} \begin{bmatrix} (\hat{w})^\top \end{bmatrix}$$

↑ $s \times 1$ vector

$$w(m,n) = \hat{w}(m) \cdot \hat{w}(n)$$

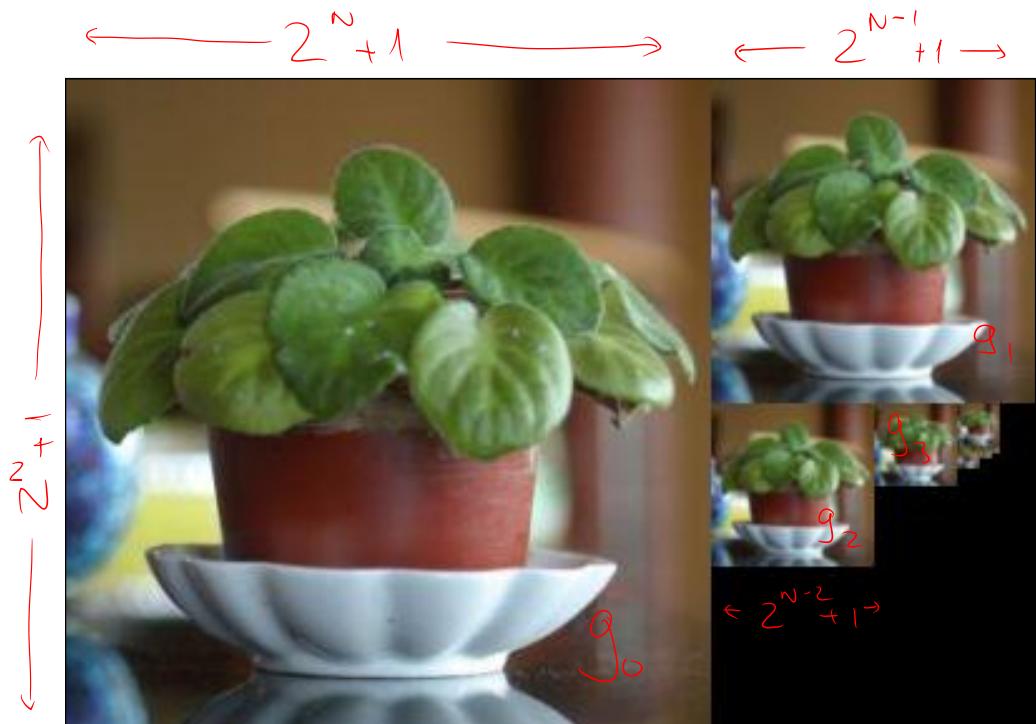


Exploiting separability to compute \hat{g}_1 :

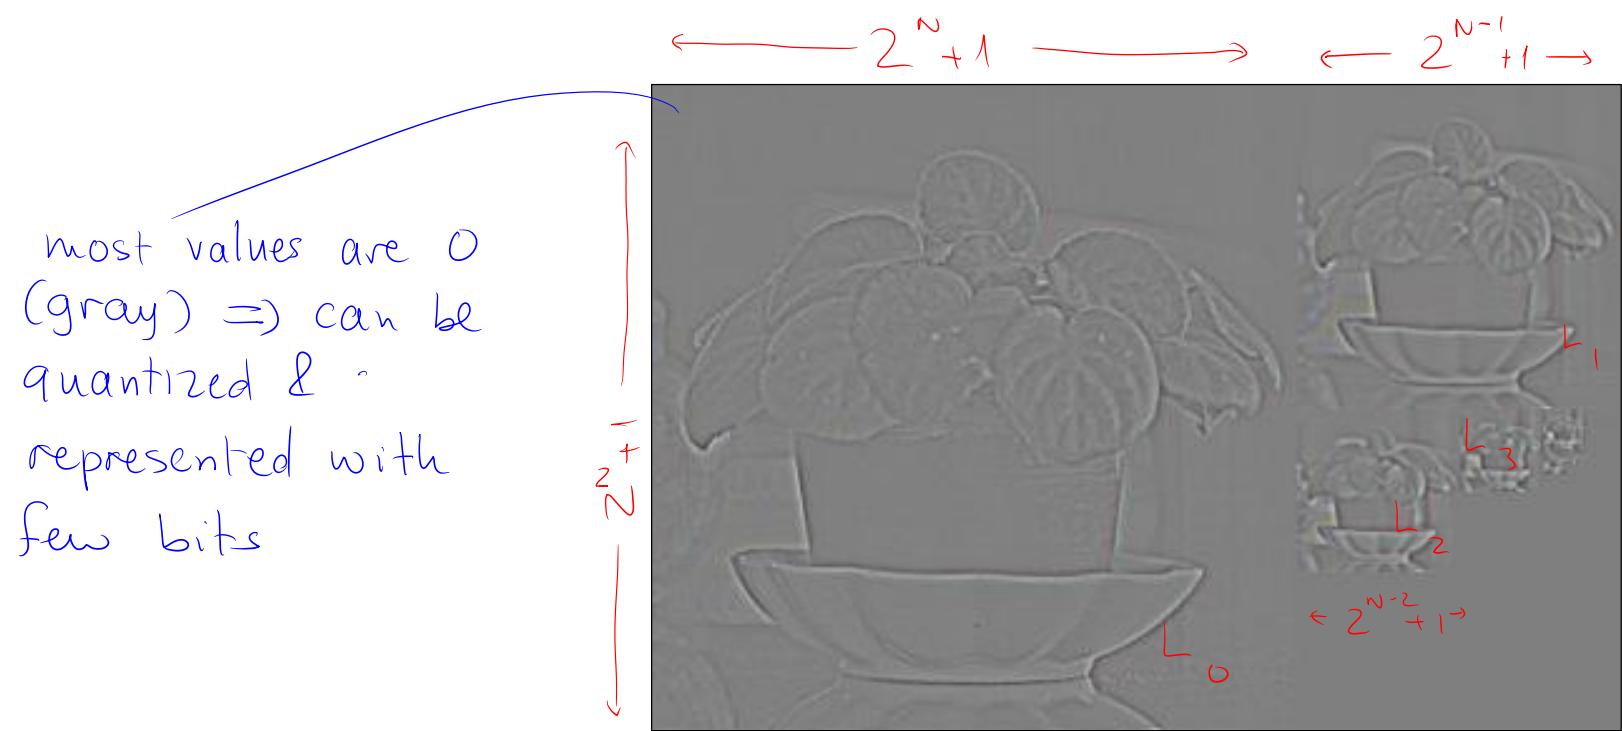
- ① Convolve each row of g_0 with \hat{w}
- ② Convolve the columns of the result with \hat{w} again

the Laplacian Pyramid

Idea: Rather than store the smoothed images, store only the difference between levels g_l and g_{l+1}



the Laplacian Pyramid



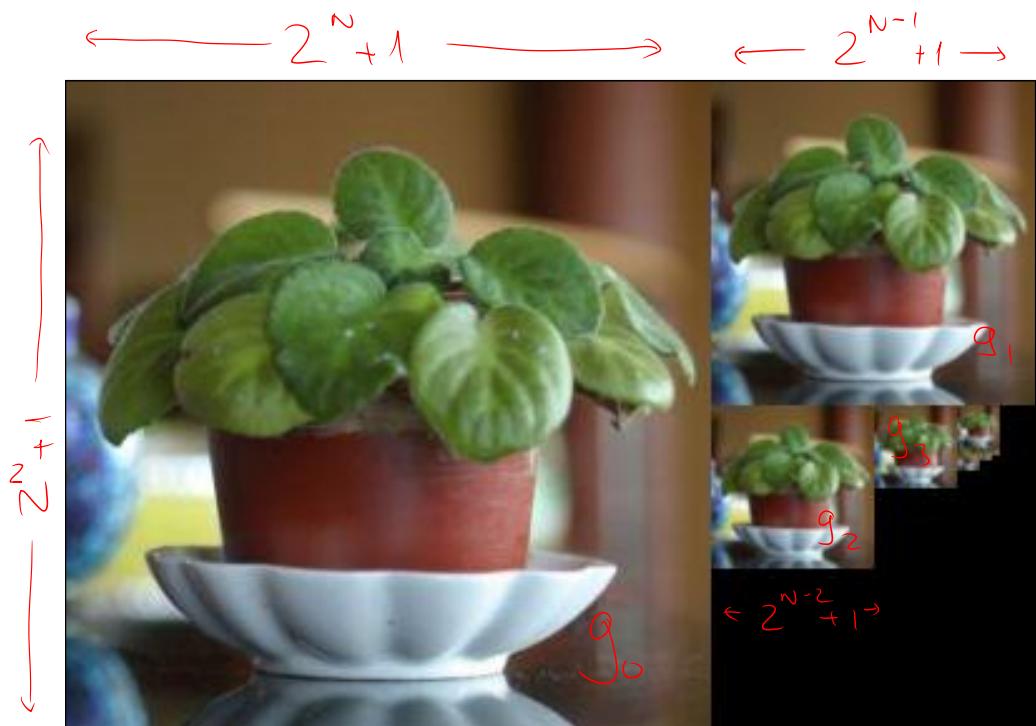
the Laplacian Pyramid

Idea: Rather than store the smoothed images, store only the difference between levels g_e and g_{e+1}

- To do this, we must compare adjacent levels g_e and g_{e+1}

- But g_{e+1}, g_e are not the same size!

\Rightarrow Expand g_{e+1} to make it equal size to g_e

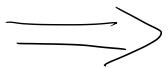


operation #3: the EXPAND() function

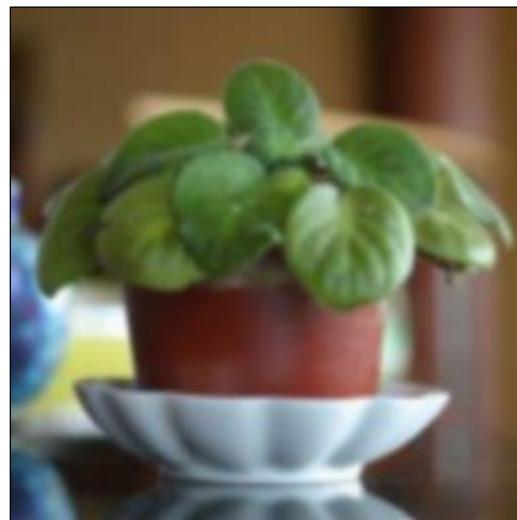
$$\leftarrow 2^{N-l} + 1 \rightarrow$$



EXPAND()



$$\leftarrow 2^{N-l+1} + 1 \rightarrow$$



g_l

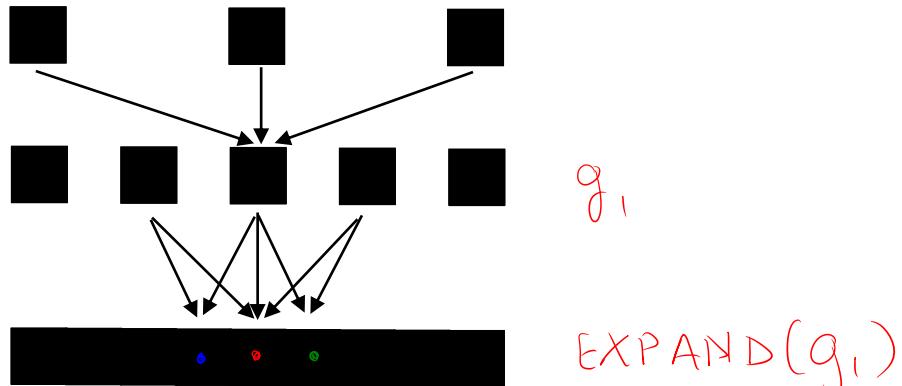
EXPAND(g_l)

$$\boxed{\text{EXPAND}(g_l) = 4 \sum_{m=-2}^2 \sum_{n=-2}^2 w(m,n) g_l\left(\frac{i-m}{2}, \frac{j-n}{2}\right)}$$

the EXPAND() function

- The function ^{upsamples} level g_e by doubling its size from $(2^{N-l}+1) \times (2^{N-l}+1)$ to $(2^{N-l+1}+1) \times (2^{N-l+1}+1)$

Visualizing in 1D



EVEN $\bullet \bullet \bullet$ ODD

pixels receive contribution

from 3 pixels above them

$$\text{(total weight} = 2c+a=\frac{1}{2})$$

from 2 pixels above them

$$\text{(total weight} = 2b=\frac{1}{2})$$

General expression in 1D

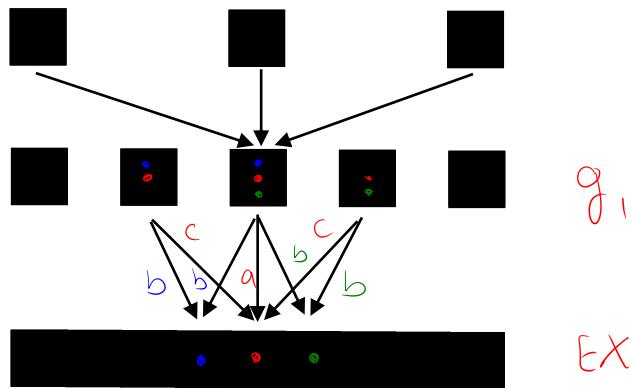
$$\text{EXPAND}(g_e)(i) = 2 \sum_{m=-2}^2 \hat{w}(m) \cdot g_e\left(\frac{i-m}{2}\right)$$

evaluated only for i, m where $\frac{i-m}{2}$ is an integer

the EXPAND() function

- The function upsamples level g_e by doubling its size from $(2^{N-l}+1) \times (2^{N-l}+1)$ to $(2^{N-l+1}+1) \times (2^{N-l+1}+1)$

Visualizing in 1D



$$\begin{aligned} & \text{EVEN pixels receive contribution from 3 pixels above them} \\ & \text{ODD pixels receive contribution from 2 pixels above them} \\ & (\text{total weight} = 2c+a=1/2) \\ & (\text{total weight} = 2b=1/2) \end{aligned}$$

EXPAND function in matrix notation (1D)

$$\begin{matrix} \uparrow & & \leftarrow 2^{N-1} \\ \text{EXPAND}(g_1) & = & \begin{bmatrix} c & b & a & b & c \\ c & b & a & b & c \\ \vdots & \ddots & \ddots & \ddots & \vdots \\ 0 & \ddots & \ddots & \ddots & \ddots \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 0 & \cdots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots \end{bmatrix} & \begin{bmatrix} g_1 \end{bmatrix} \end{matrix}$$

Convolution matrix $(D_0)^T$
upsampling matrix C_0

$$\boxed{\text{EXPAND}(g_1) = C_0 \cdot (D_0)^T g_1}$$

the Laplacian Pyramid



the Laplacian Pyramid

$$L_1 = g_1 - \text{EXPAND}(g_2)$$



g_2



L_1

the Laplacian Pyramid

$$L_2 = g_2 - \text{EXPAND}(g_3)$$



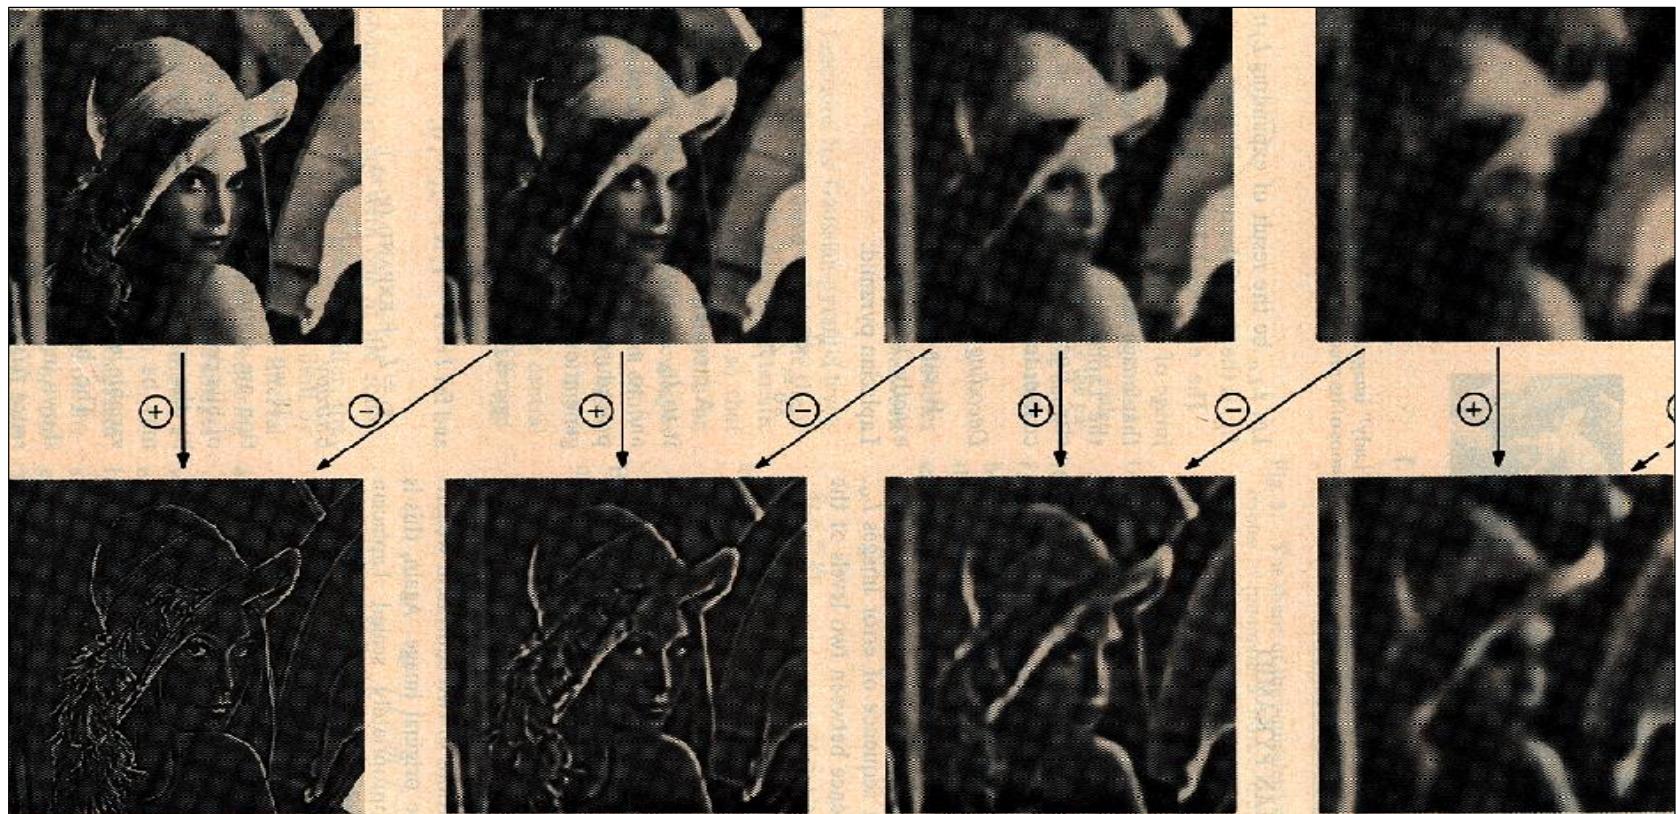
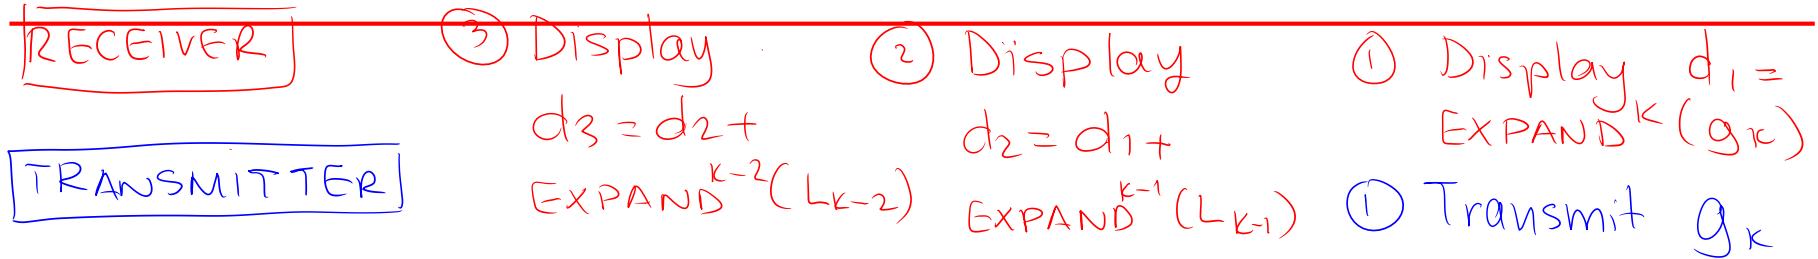
the Laplacian Pyramid

Often we use a truncated Laplacian pyramid by storing images $l_0, l_1, \dots, l_k, g_{k+1}$ for $k+1 < N$

Base case: store g_N



application: progressive transmission



⑤ Transmit
 L_{K-4}

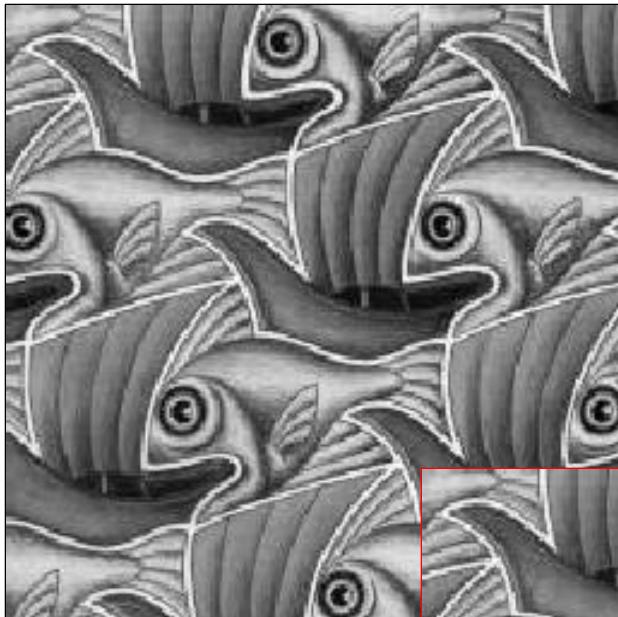
④ Transmit
 L_{K-3}

③ Transmit
 L_{K-2}

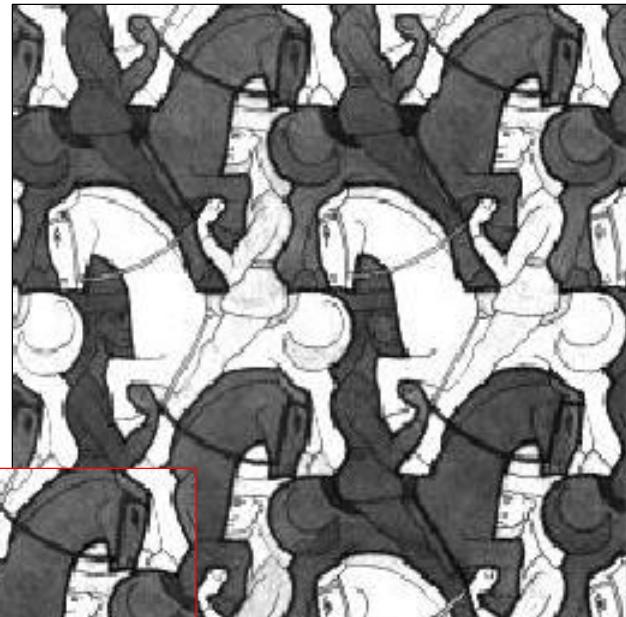
② Transmit
 L_{K-1}

application: Laplacian image blending

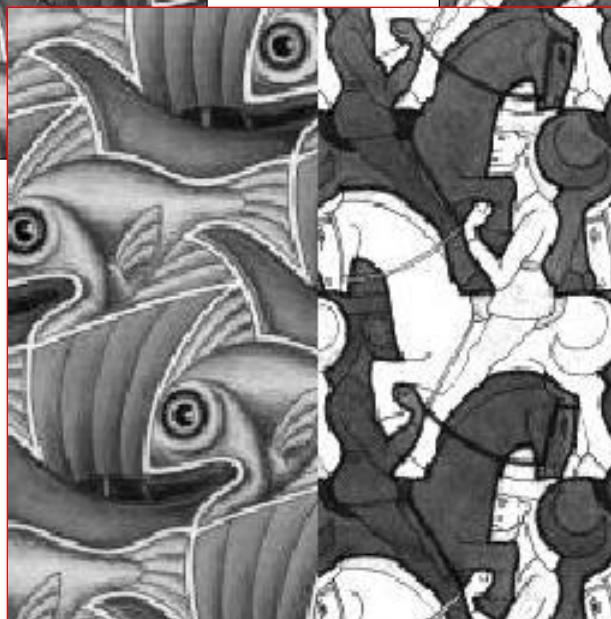
Source 1



Source 2



Blend

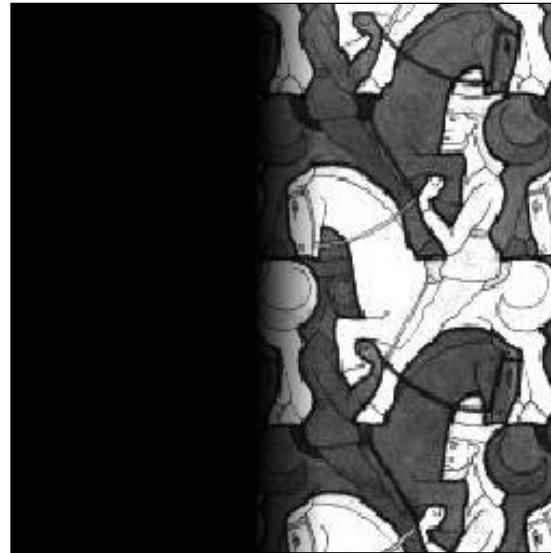


Slides adapted
from A. Efros
(CMU)

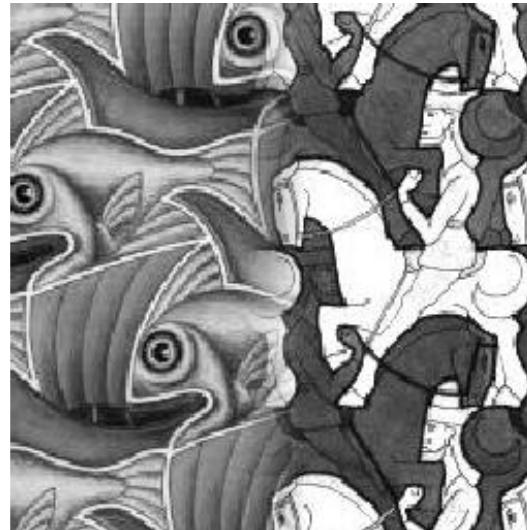
feathering



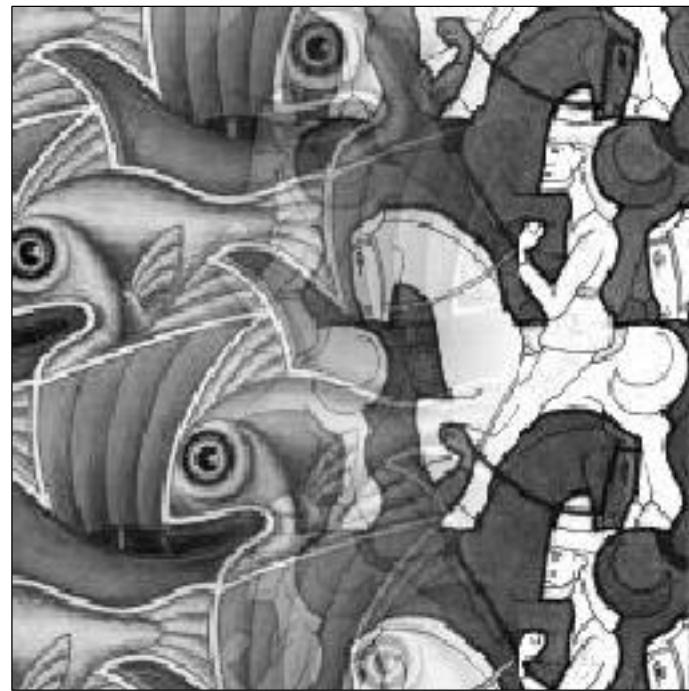
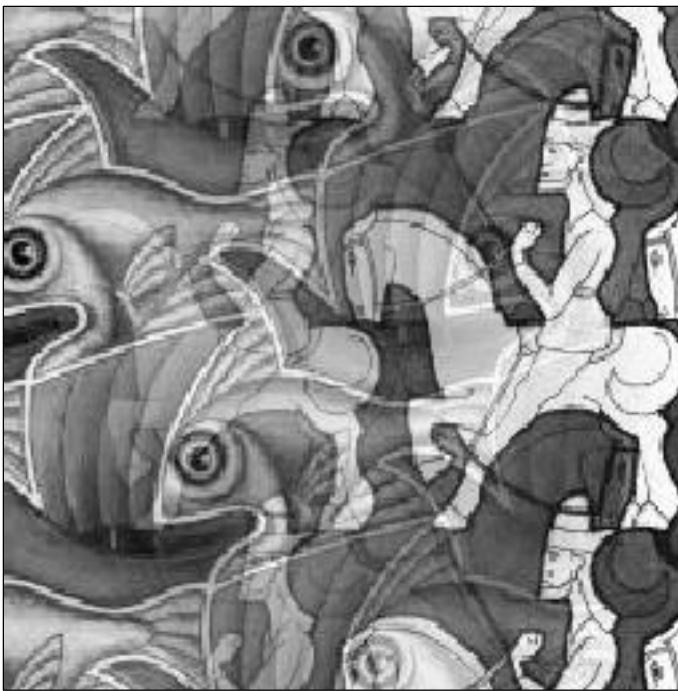
+



-

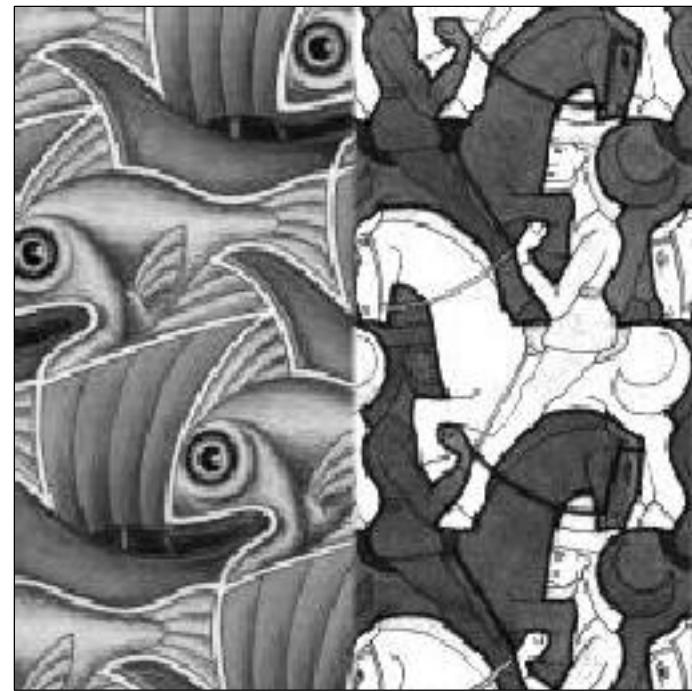
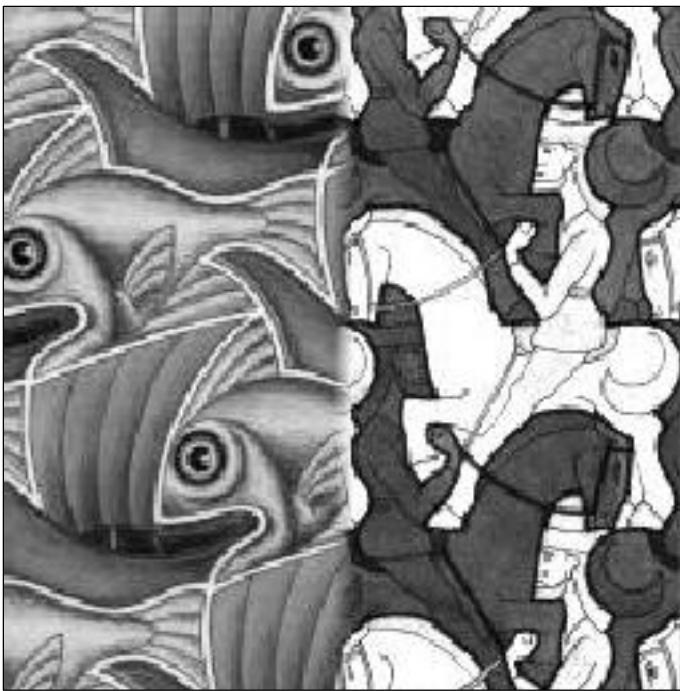


feathering: effect of Window Size



Ghosting is most
apparent here

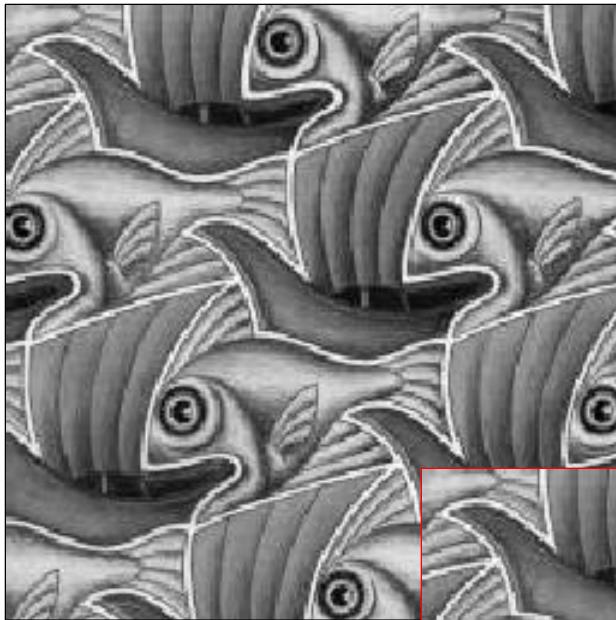
feathering: effect of Window Size



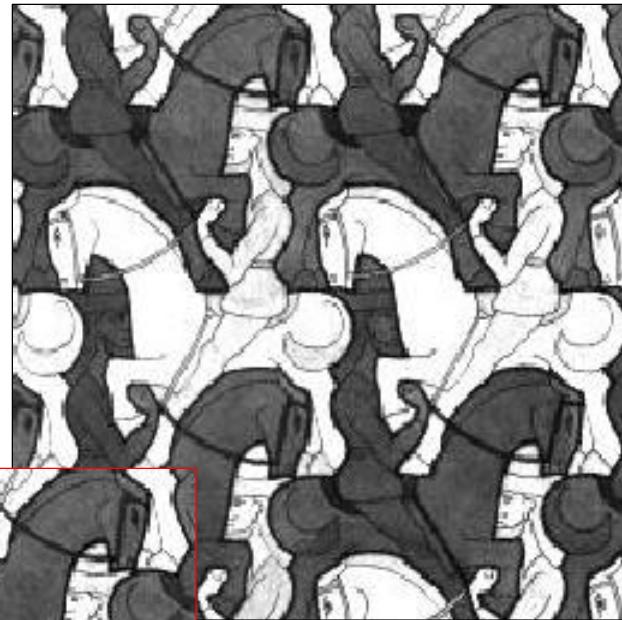
seam most visible
here

Laplacian image blending (Burt & Adelson, 1983)

Source A



Source B



Blend

Main challenge:

- Minimize ghosting without making visible seams



Slides adapted
from A. Efros
(CMU)

pyramid blending algorithm

Input: Source images A, B & binary matte M
Output: Blended image S

A

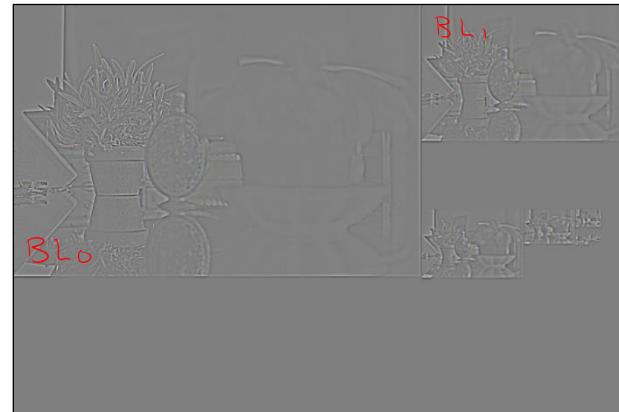
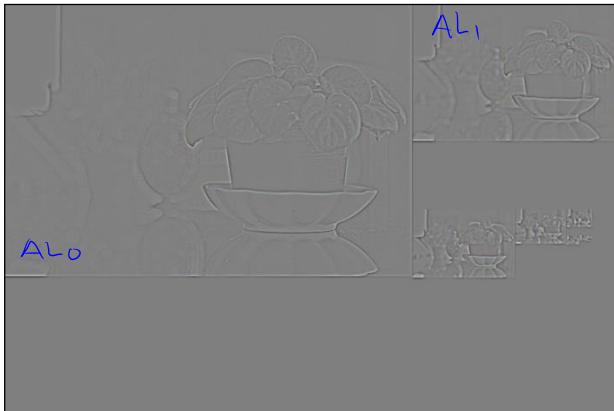


B



pyramid blending algorithm

Iput: Source images A, B & binary matte M
Output: Blended image S



① Compute A's
Laplacian
pyramid:
 $AL_0, \dots, AL_{N-1}, AG_N$

③ Compute M's
Gaussian
pyramid:
 Mg_0, \dots, Mg_N

② Compute B's
Laplacian
pyramid:
 $BL_0, \dots, BL_{N-1}, BG_N$

pyramid blending algorithm

④ Compute the Laplacian pyramid, $SL_0, SL_1, \dots, SL_{N-1}, Sg_N$ by applying the matting equation with matte Mg_e :

$$SL_\ell(i,j) = Mg_\ell(i,j) AL_\ell(i,j) + (1 - Mg_\ell(i,j)) BL_\ell(i,j)$$

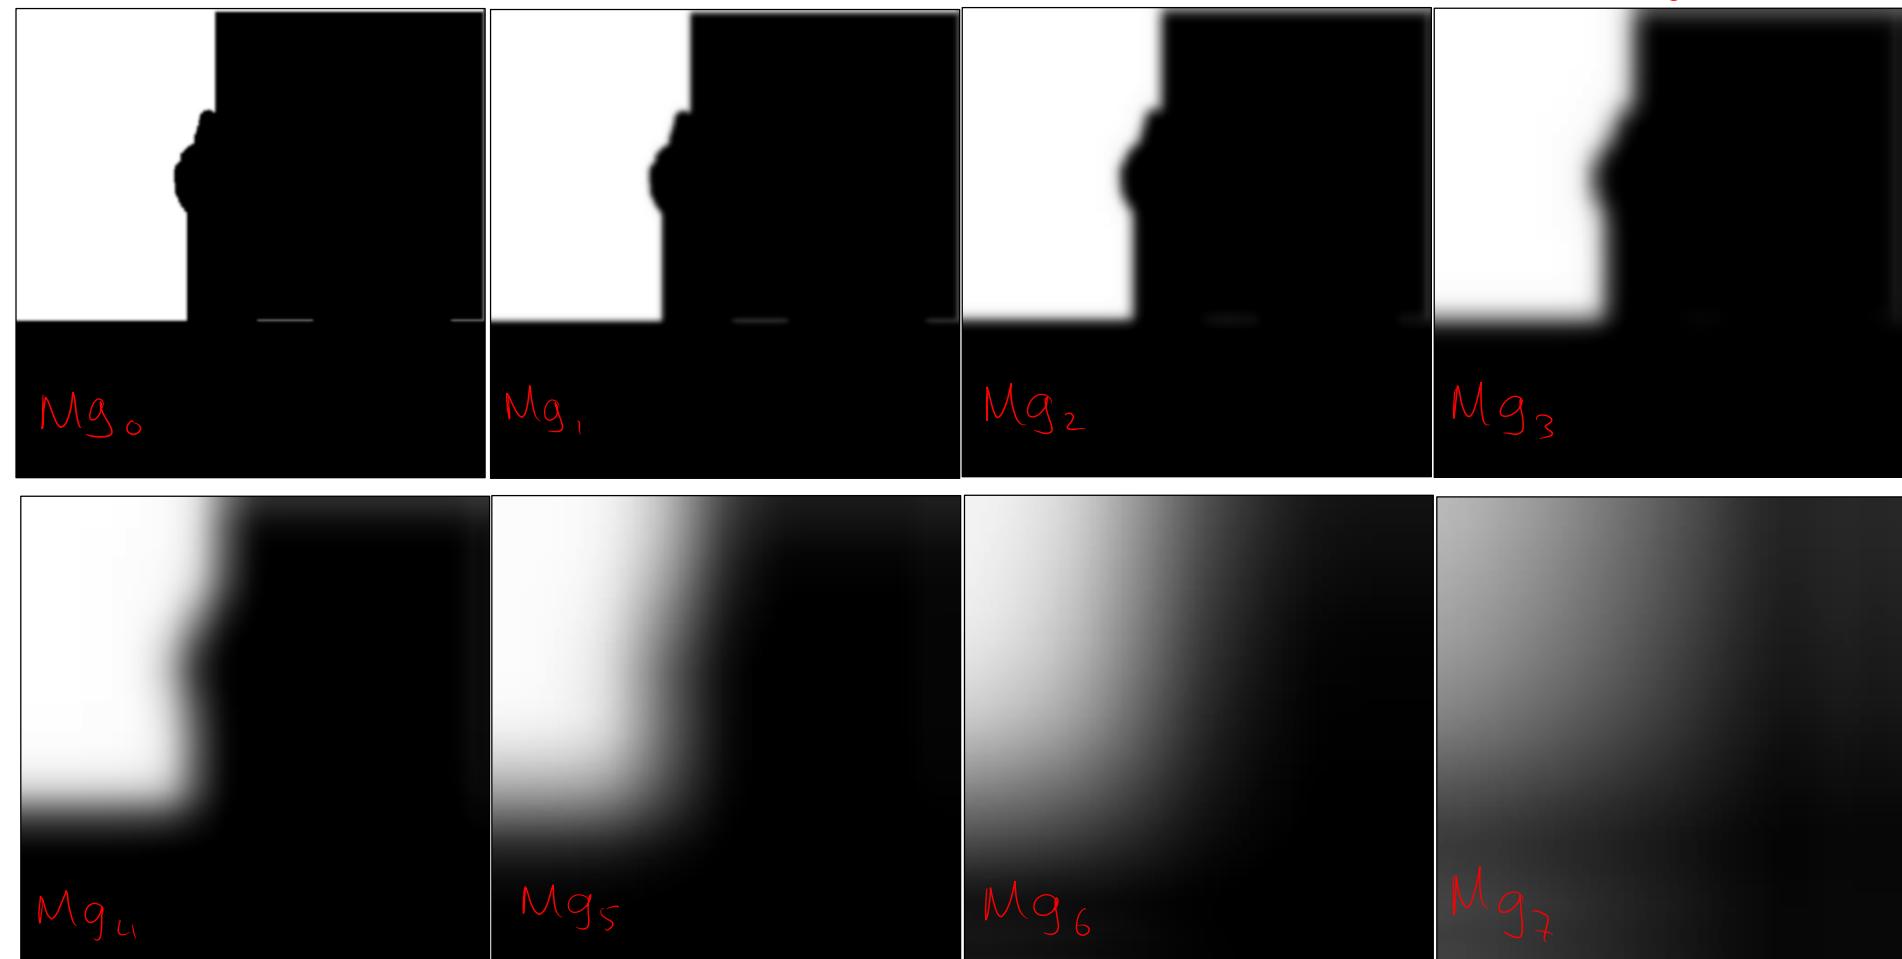


pyramid blending algorithm: intuition

The algorithm effectively uses a different alpha matte for each level of detail

\uparrow detail \Rightarrow \downarrow feathering window

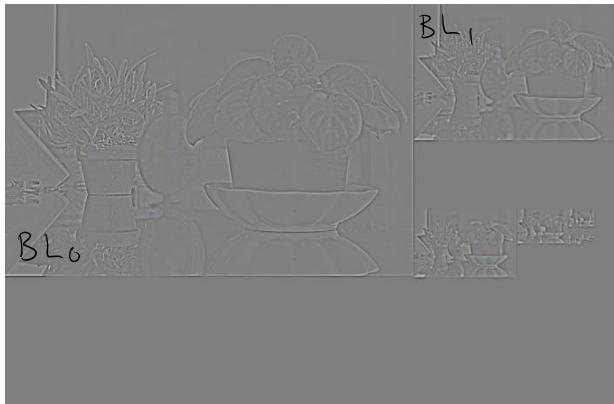
\downarrow detail \Rightarrow \uparrow feathering window



pyramid blending algorithm

⑤ Compute level 0 of the Gaussian pyramid of S from $S_{L0}, \dots, S_{Ln-1}, S_{gn}$

Result S'



stitching photos with pyramid blending



stitching photos (no pyramid blend)



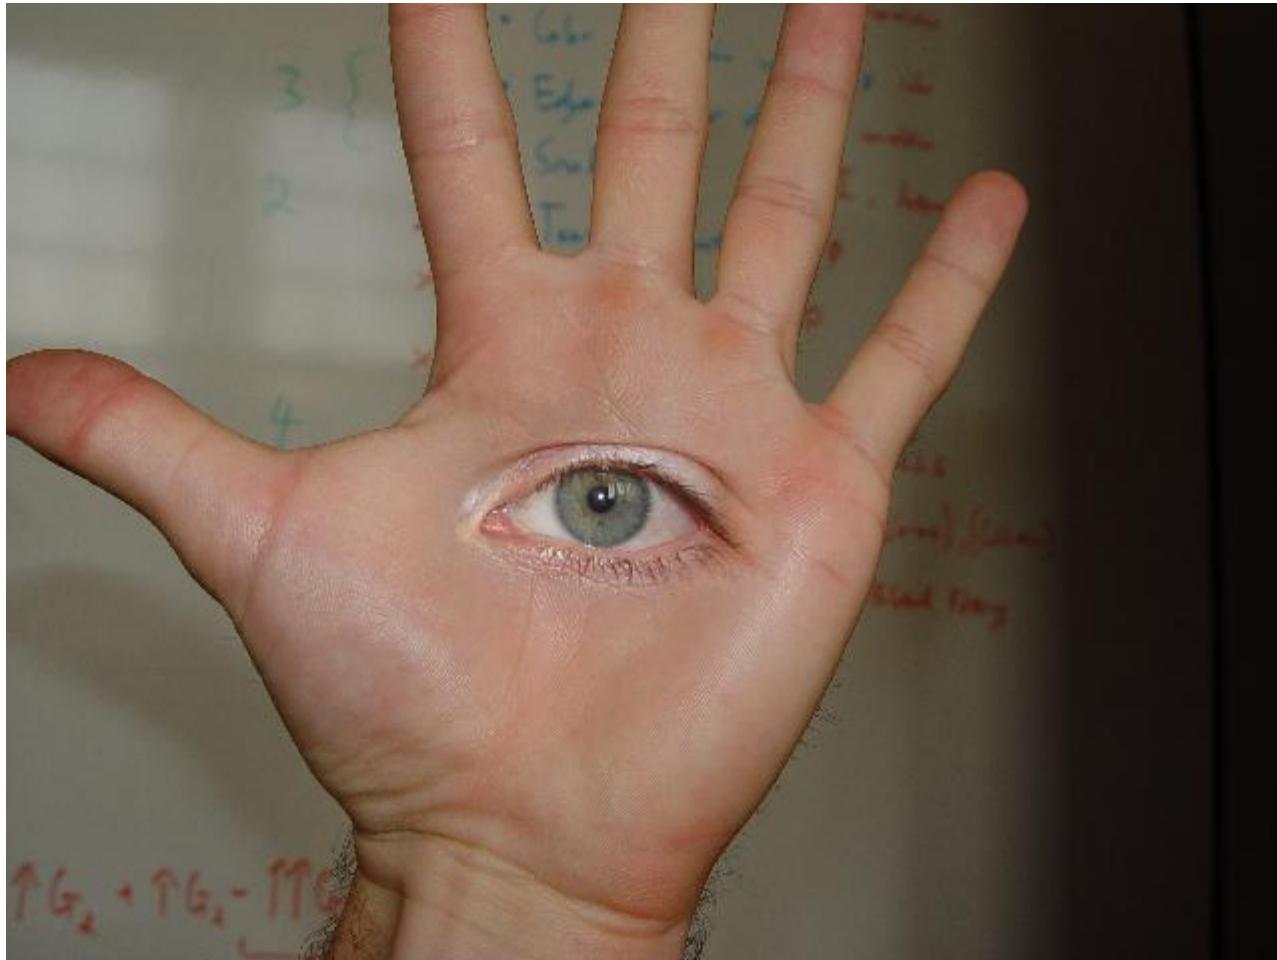
blending mismatched photos (still looks OK)



stitching mismatched photos (no pyramid blend)



Horror Photo



© prof. d'martin

Topic 07:

Multi-scale image transforms

- Gaussian & Laplacian pyramids
- Linear image transforms
- Haar wavelet transform

linear transform framework

Projection Vectors: Let $\vec{\mathbf{I}}$ denote a 1D signal, or a vectorized representation of an image (so $\vec{\mathbf{I}} \in \mathcal{R}^N$), and let the transform be

$$\vec{\mathbf{a}} = \mathbf{P}^T \vec{\mathbf{I}}. \quad (1)$$

Here,

- $\vec{\mathbf{a}} = [a_0, \dots, a_{M-1}] \in \mathcal{R}^M$ are the transform coefficients.
- The columns of $\mathbf{P} = [\vec{\mathbf{p}}_0, \vec{\mathbf{p}}_1, \dots, \vec{\mathbf{p}}_{M-1}]$ are the projection vectors; the m^{th} coefficient, a_m , is the inner product $\vec{\mathbf{p}}_m^T \vec{\mathbf{I}}$
- When \mathbf{P} is complex-valued, we should replace \mathbf{P}^T by the conjugate transpose \mathbf{P}^{*T}

Sampling: The transform $\mathbf{P}^T \in \mathcal{R}^{M \times N}$ is said to be *critically sampled* when $M = N$. Otherwise it is *over-sampled* when $M > N$, or *under-sampled* when $M < N$.

Basis Vectors: For many transforms of interest there is a corresponding basis matrix \mathbf{B} satisfying

$$\vec{\mathbf{I}} = \mathbf{B} \vec{\mathbf{a}}. \quad (2)$$

linear transform framework

Completeness

- the forward transform (1) is complete, encoding all image structure, if it is invertible.
- when critically sampled, it is complete if $\mathbf{B} = (\mathbf{P}^T)^{-1}$ exists.
- if over-sampled, the transform is complete if $\text{rank}(\mathbf{P}) = N$.
In this case \mathbf{B} is not unique – one choice is the pseudoinverse

$$\mathbf{B} = (\mathbf{P}^T \mathbf{P})^{-1} \mathbf{P}^T$$

- if under-sampled, then $\text{rank}(\mathbf{P}) < N$ and it is not invertible in general.

Self-Inverting

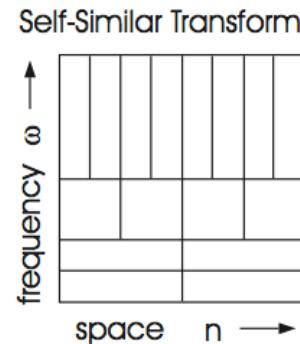
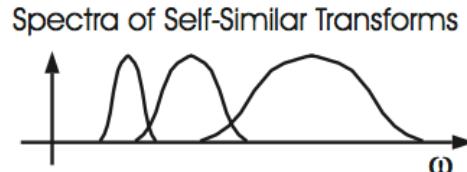
- the transform is self-inverting if $\vec{\mathbf{b}}_m = \alpha \vec{\mathbf{p}}_m$ for some constant α .
- in the critically-sampled, self-inverting case the transform is orthogonal (unitary), up to the constant α (e.g., the DFT).

self-similar linear transforms

Goal: The filter support should grow with scale, and be well matched to scale-dependent correlation lengths in images. The representation should exhibit scale-invariant properties, as objects project to images at different scales depending on distance from camera.

Scale Self-Similarity: Let the basis functions be dilations and translations of a “mother” function, so they all have the same shape, differing in scale and position only.

Gabor Wavelet
Basis Functions



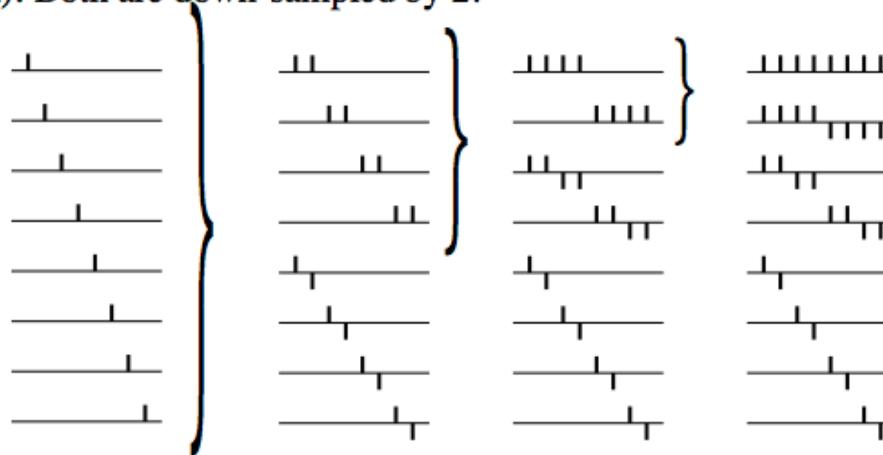
Topic 07:

Multi-scale image transforms

- Gaussian & Laplacian pyramids
- Linear image transforms
- Haar wavelet transform

Haar wavelet transform

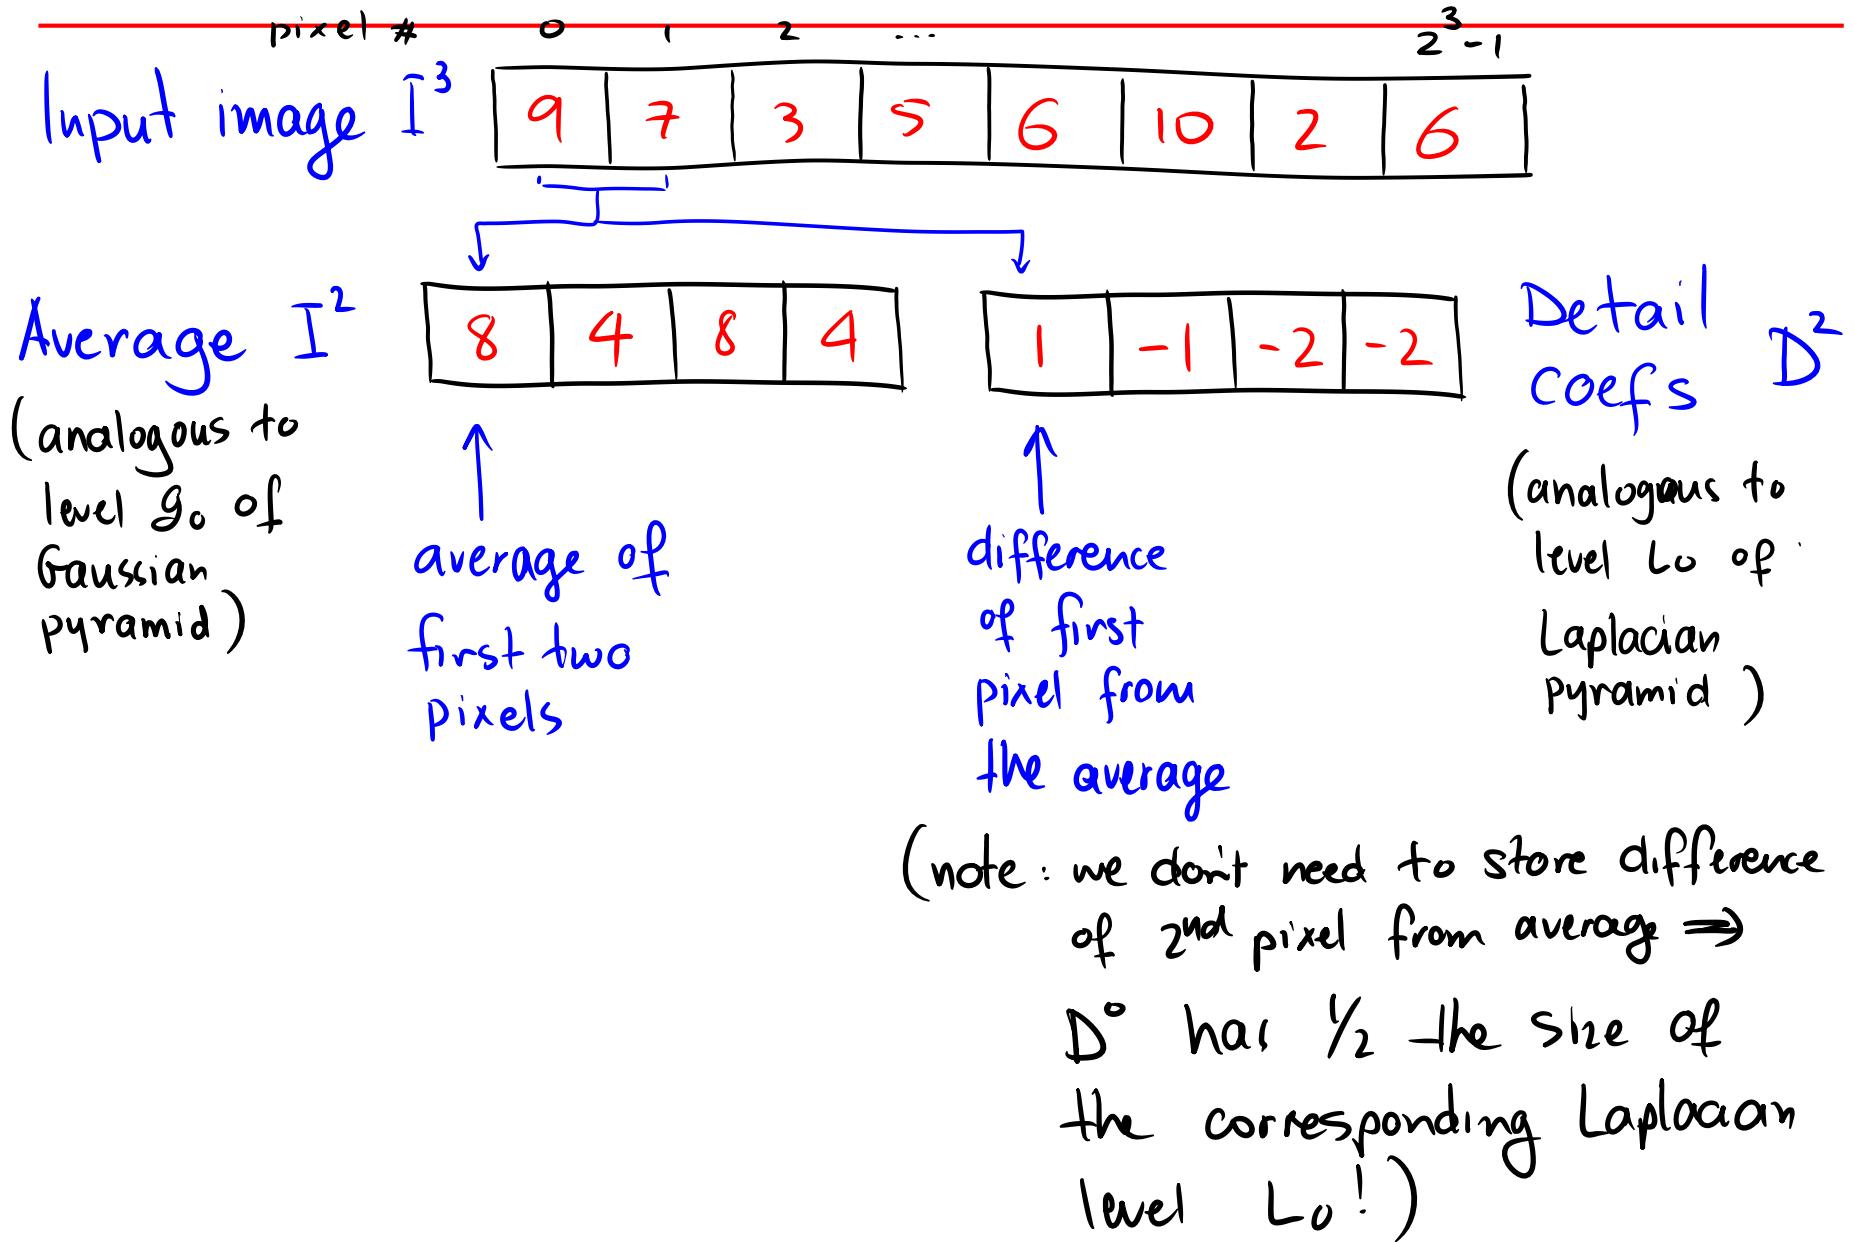
Originally described by A. Haar (1909). Each step creates two channels: one simply averages adjacent elements (i.e., low-pass channel); and one takes difference between adjacent elements (i.e., a high-pass channel). Both are down-sampled by 2.



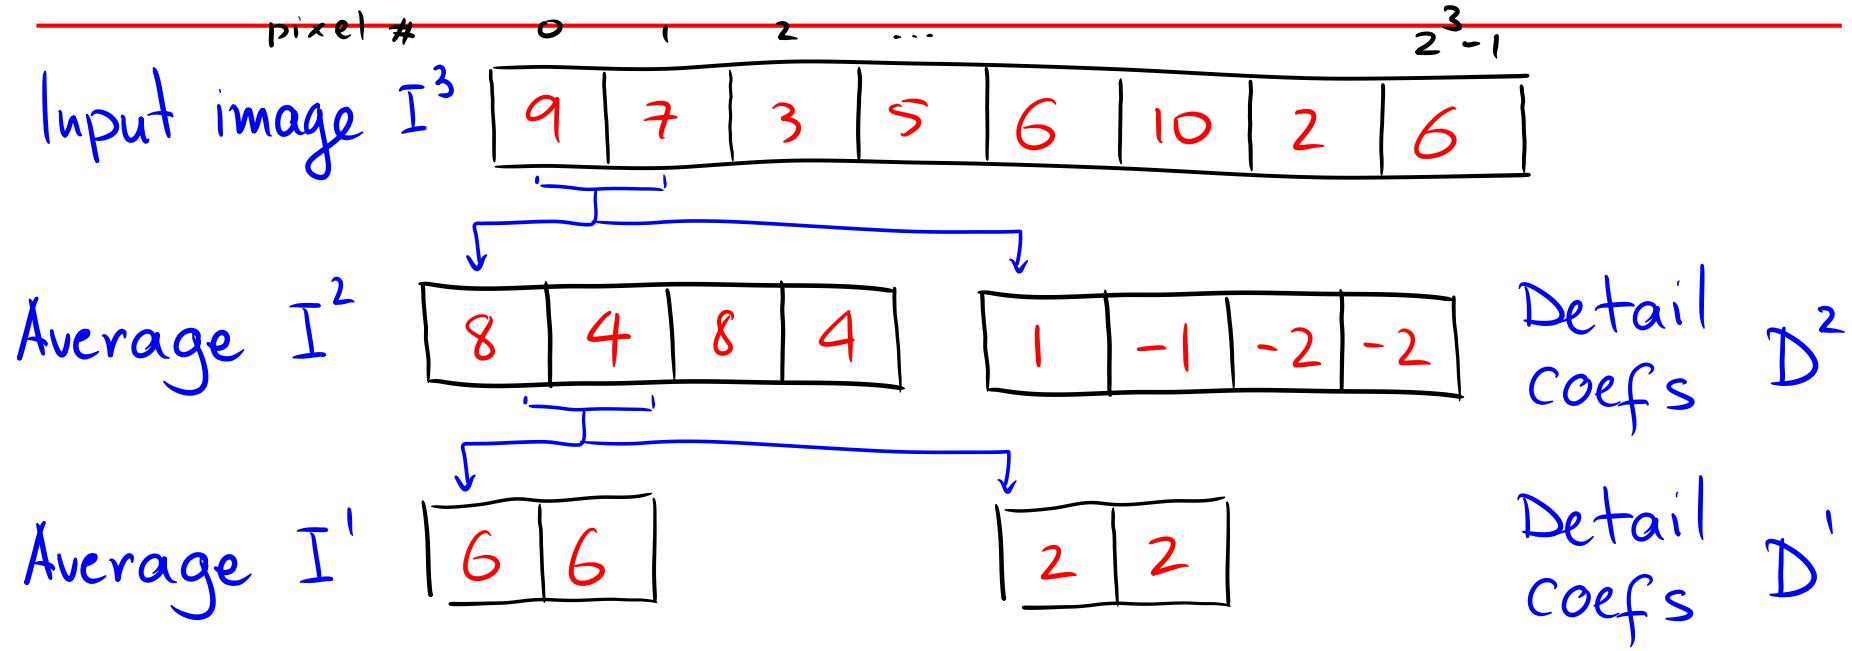
Properties:

- critically-sampled and self-inverting (orthogonal)
- local in space (compact) but not continuously differentiable
- broad ringing frequency spectrum due to top-hat spatial window, and therefore massive aliasing in each band (like blocked DCT).
- very efficient to compute with pyramid scheme and addition

1D Haar wavelet transform: recursive definition



1D Haar wavelet transform: recursive definition



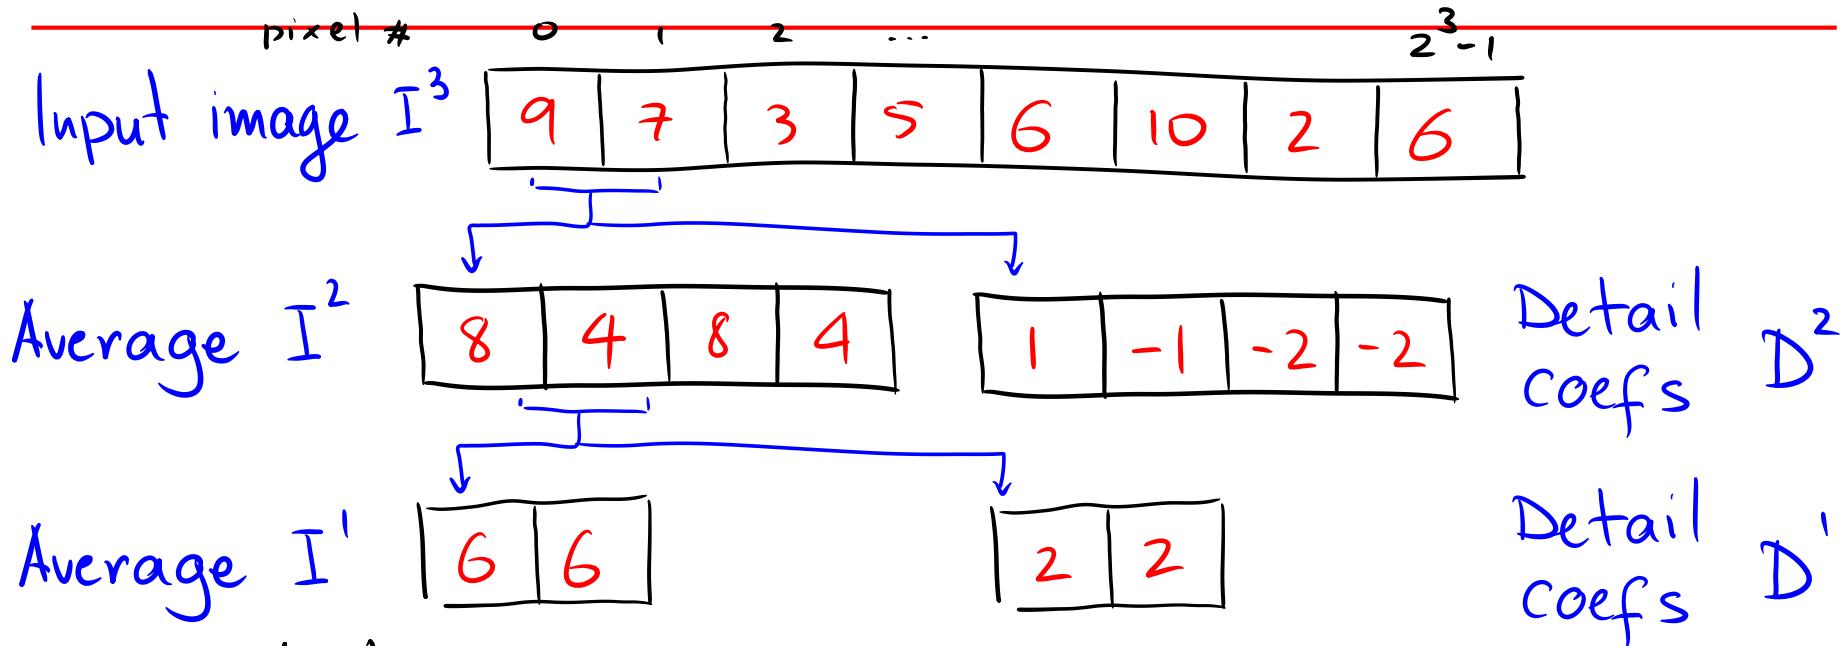
$$I_i^j = \frac{1}{2} (I_{2i}^{j+1} + I_{2i+1}^{j+1})$$

j -th level of "pyramid" contains 2^j pixels

$$D_i^j = I_{2i}^{j+1} - \frac{1}{2} (I_{2i}^{j+1} + I_{2i+1}^{j+1})$$

$$= \frac{1}{2} (I_{2i}^{j+1} - I_{2i+1}^{j+1})$$

1D Haar wavelet transform: recursive definition



equivalent
convolution
mask ϕ :

$$\begin{bmatrix} 1/2 & 1/2 \end{bmatrix}$$

equivalent
convolution
mask ψ :

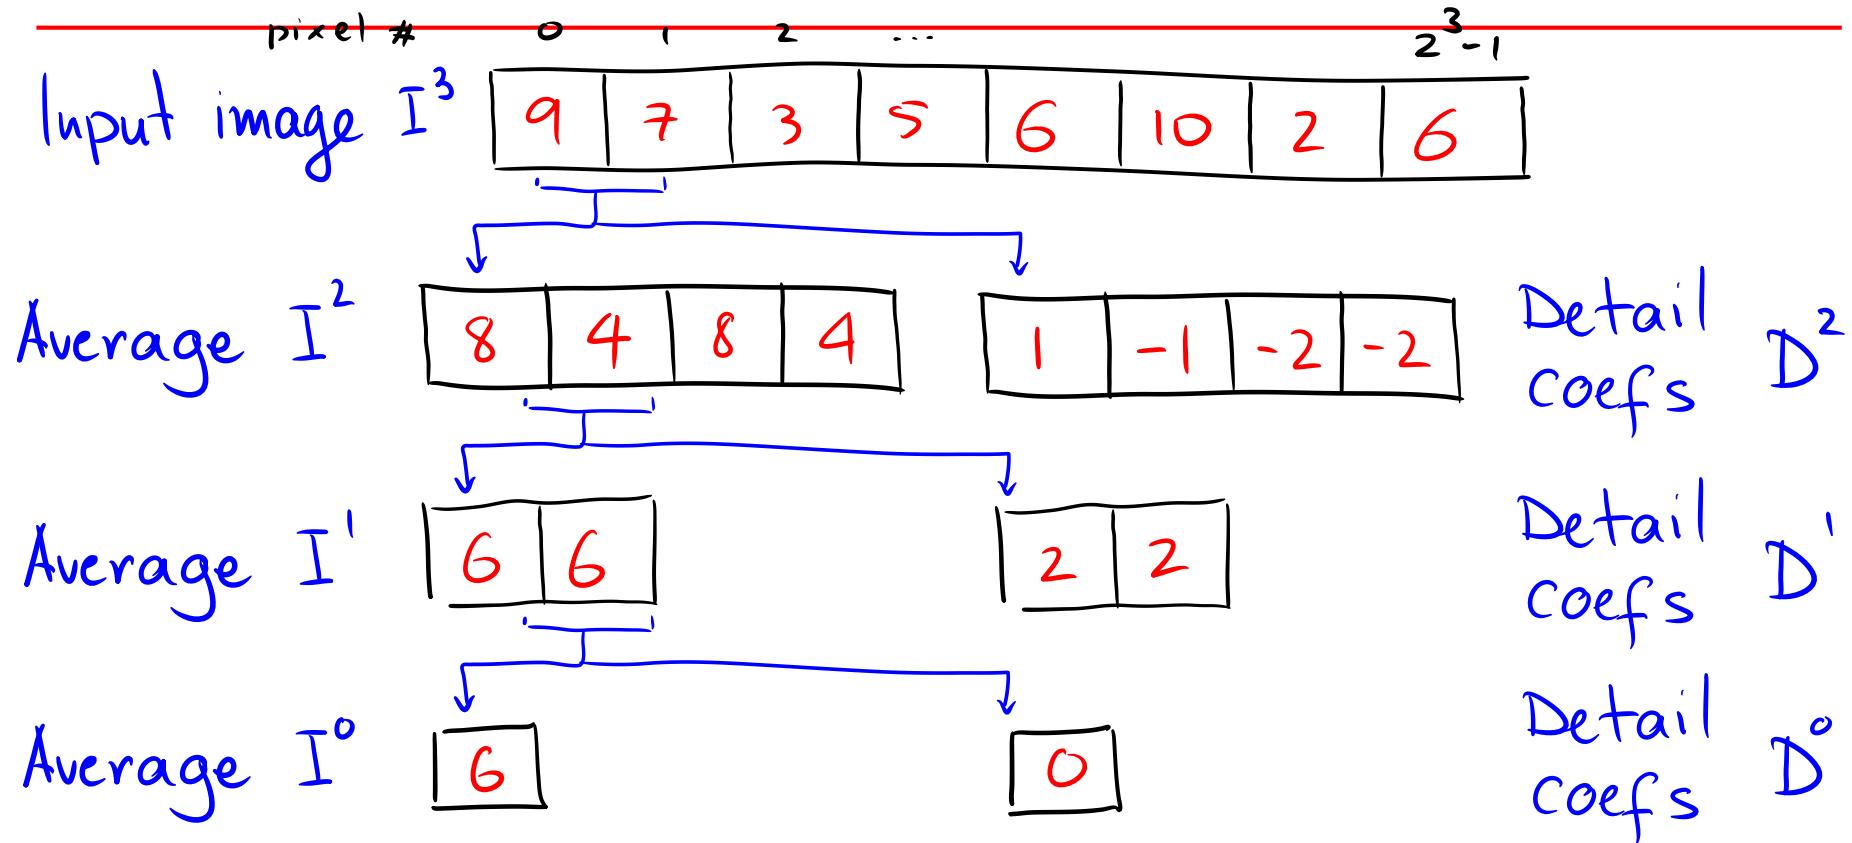
$$\begin{bmatrix} -1/2 & 1/2 \end{bmatrix}$$

$$I_i^j = \frac{1}{2} (I_{2i}^{j+1} + I_{2i+1}^{j+1})$$

$$D_i^j = \frac{1}{2} (I_{2i}^{j+1} - I_{2i+1}^{j+1})$$

j-th level of "pyramid" contains
 2^j pixels

1D Haar wavelet transform: recursive definition

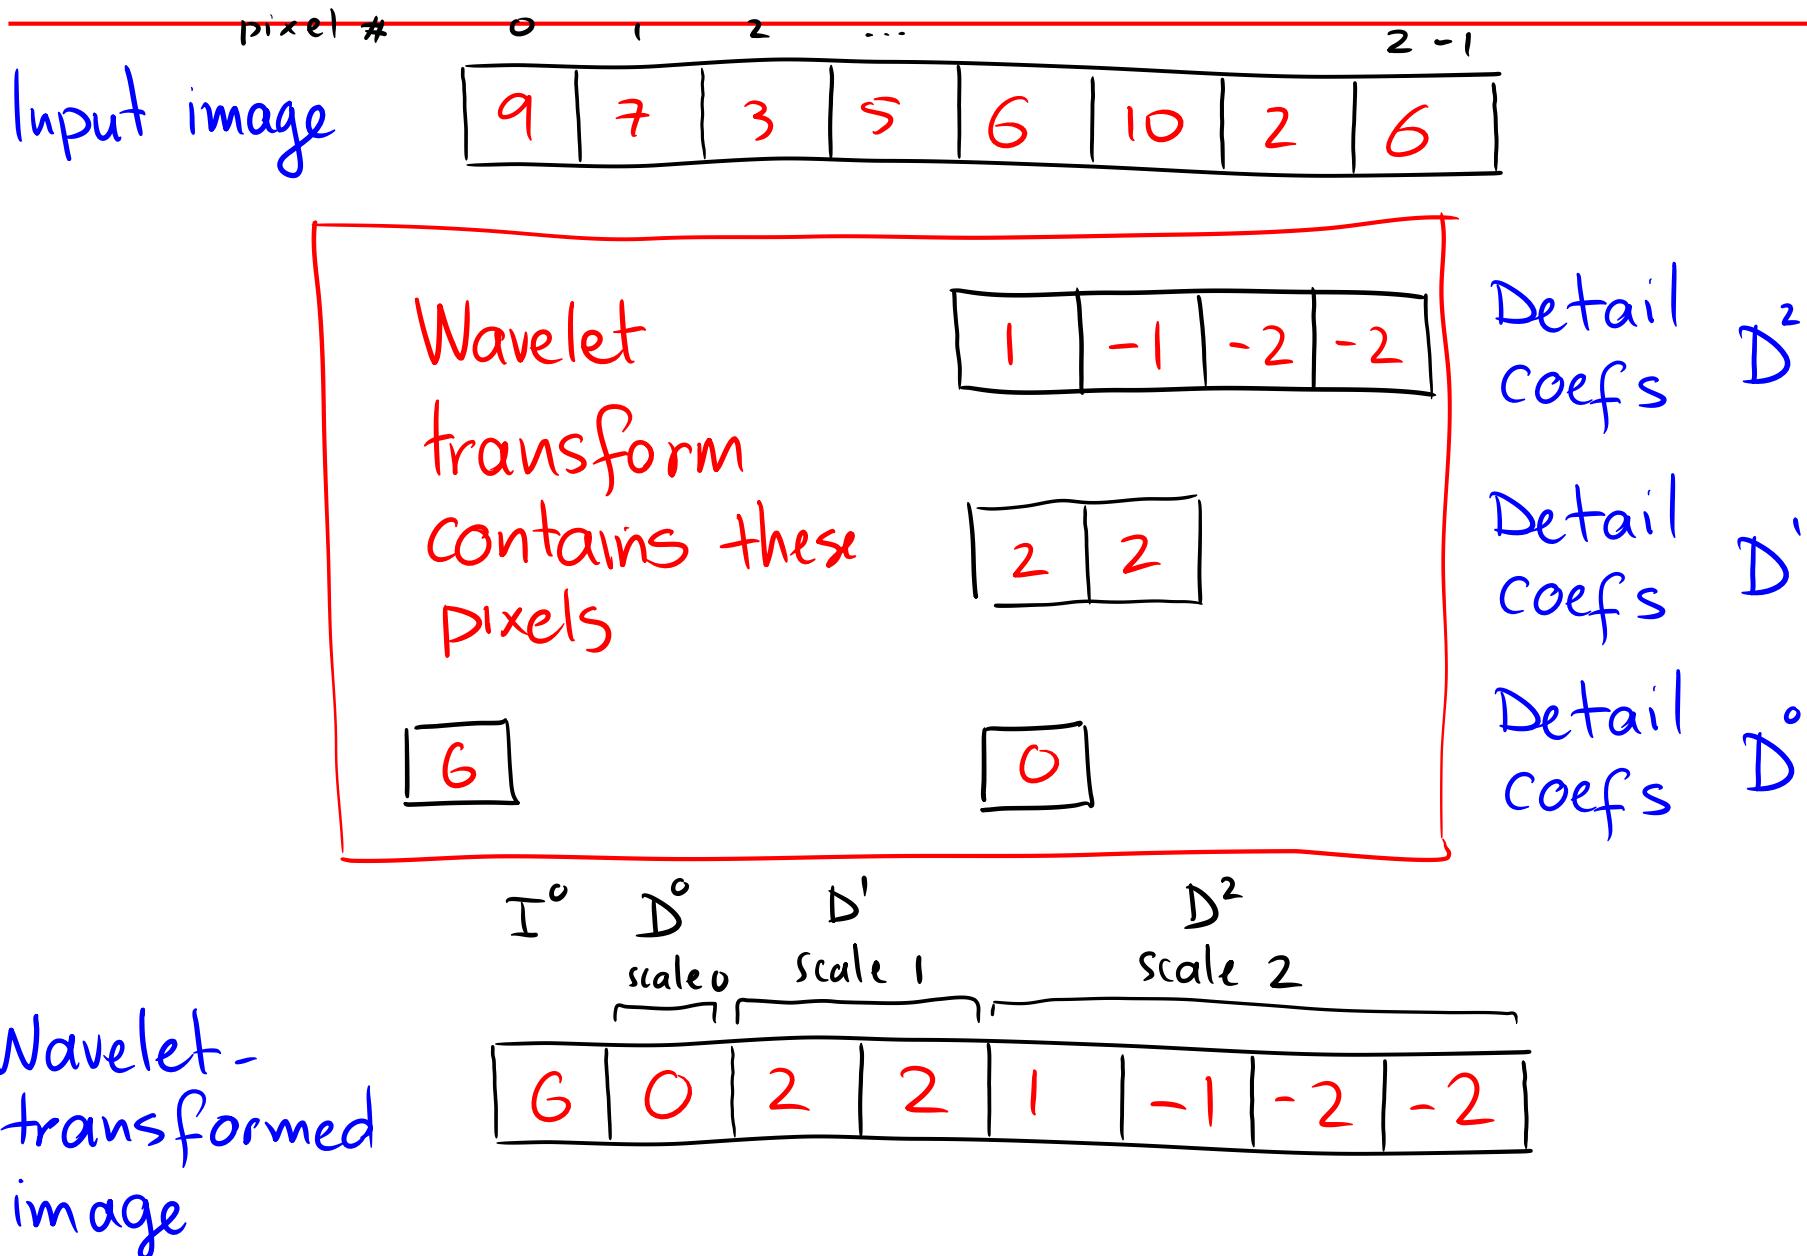


$$I_i^j = \frac{1}{2} (I_{2i}^{j+1} + I_{2i+1}^{j+1})$$

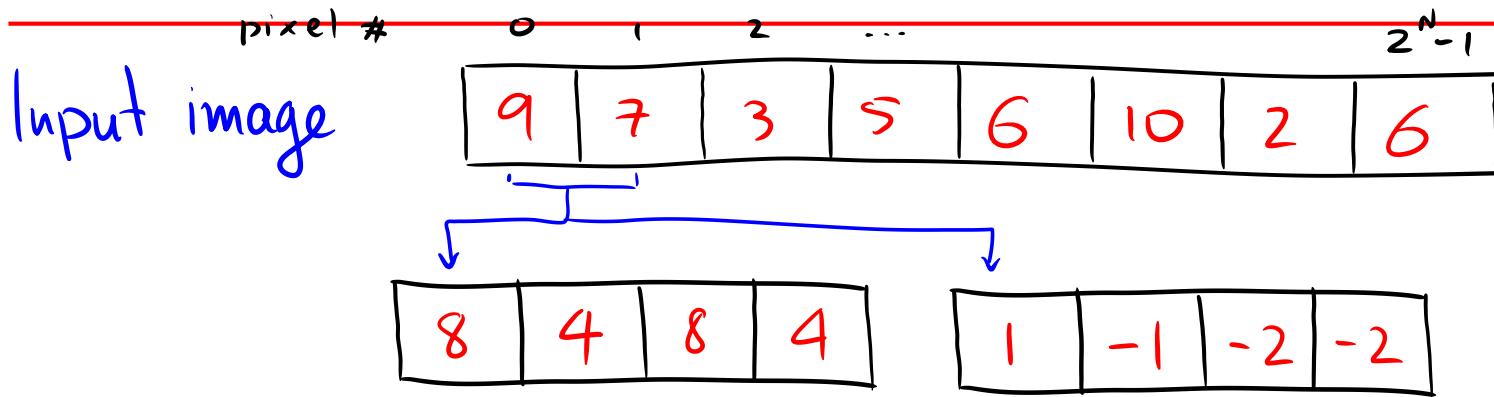
$$D^j = \frac{1}{2} (I_{2i}^{j+1} - I_{2i+1}^{j+1})$$

j-th level of "pyramid" contains 2^j pixels

1D Haar wavelet transform: recursive definition



1D Haar wavelet transform: matrix formulation

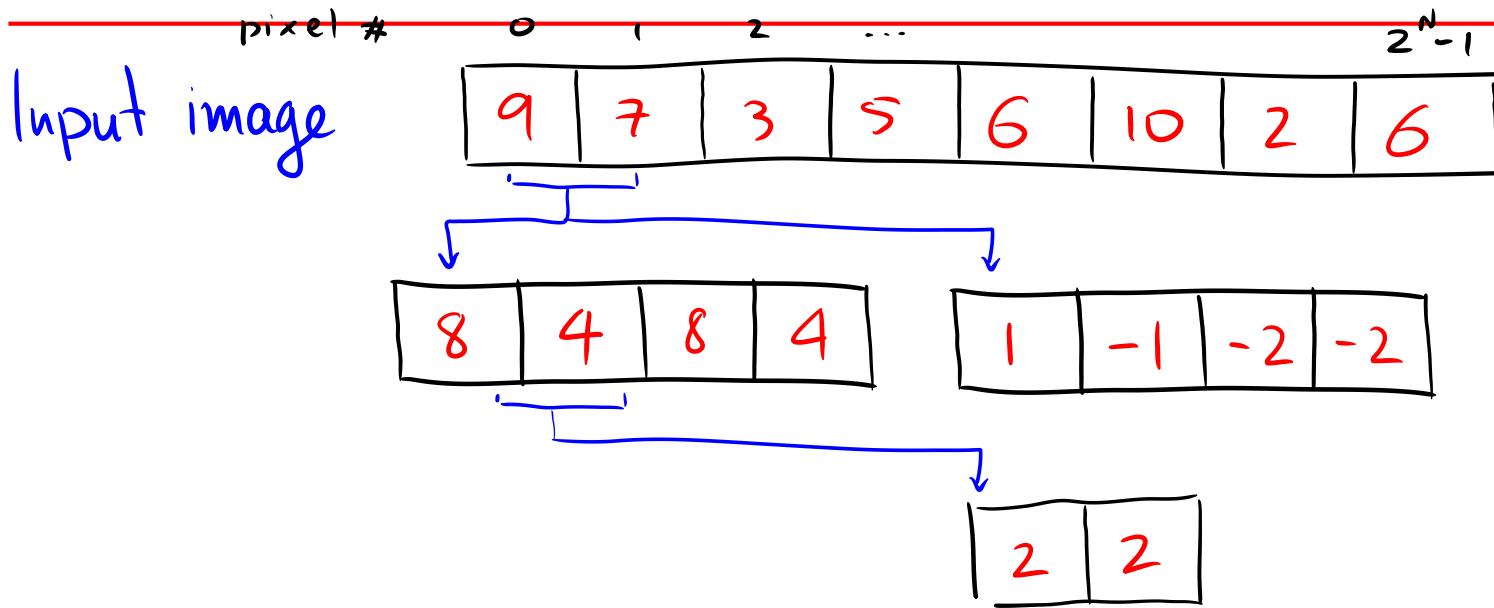


Wavelet transformed image

$$\begin{bmatrix} I^0 \\ D^0 \\ D^1 \\ D^2 \end{bmatrix} = \begin{bmatrix} 6 \\ 0 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} 9 \\ 7 \\ 3 \\ 5 \\ 6 \\ 10 \\ 2 \\ 6 \end{bmatrix}$$

Original image

1D Haar wavelet transform: matrix formulation

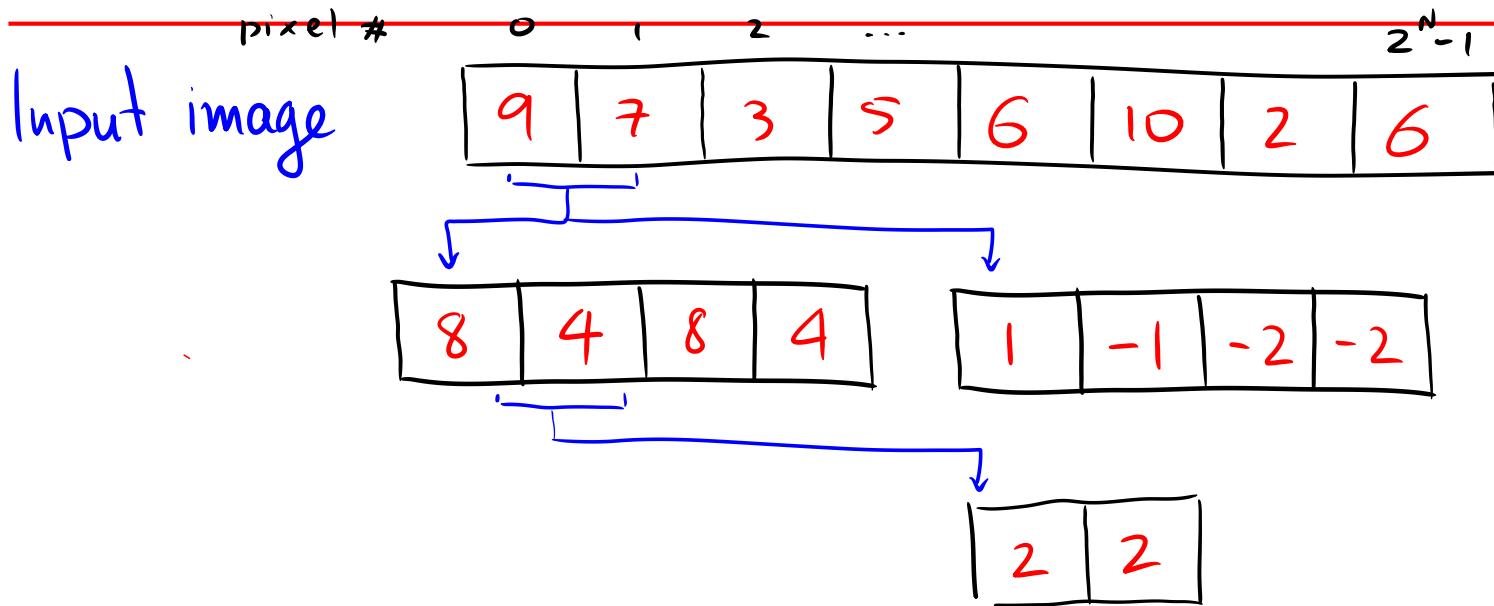


Wavelet transformed image

$$\begin{bmatrix} I^0 \\ D^0 \\ D' \\ D^2 \end{bmatrix} = \frac{1}{\sqrt{4}} \begin{bmatrix} 1 & 1 & -1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & -1 & -1 \\ 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} 9 \\ 7 \\ 3 \\ 5 \\ 6 \\ 10 \\ 2 \\ 6 \end{bmatrix}$$

Original image

1D Haar wavelet transform: matrix formulation



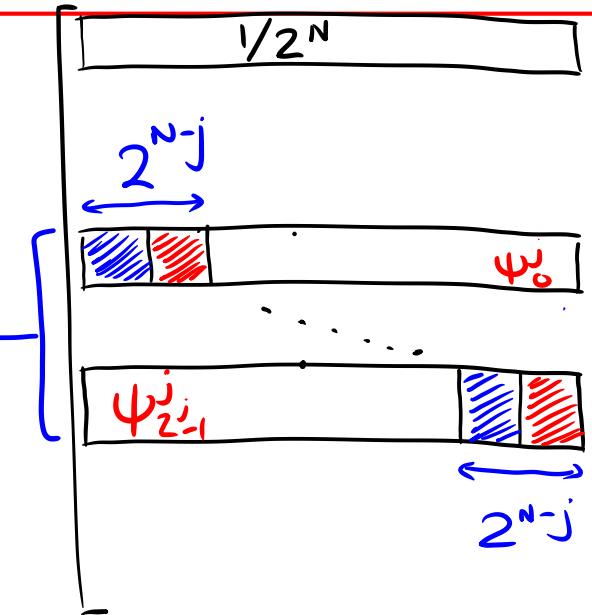
Wavelet transformed image

$$\begin{matrix}
 I^0 \\
 D^0 \\
 D' \\
 D^2
 \end{matrix} = \begin{matrix}
 \frac{1}{\sqrt{2}} & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\
 \frac{1}{\sqrt{2}} & 1 & 1 & 1 & 1 & -1 & -1 & -1 \\
 \frac{1}{\sqrt{4}} & 1 & 1 & -1 & -1 & 0 & 0 & 0 \\
 \frac{1}{\sqrt{8}} & 0 & 0 & 0 & 0 & 1 & 1 & -1 \\
 \frac{1}{\sqrt{8}} & 1 & -1 & 0 & 0 & 0 & 0 & 0 \\
 \frac{1}{\sqrt{16}} & 0 & 0 & 1 & -1 & 0 & 0 & 0 \\
 \frac{1}{\sqrt{16}} & 0 & 0 & 0 & 0 & 1 & -1 & 0 \\
 \frac{1}{\sqrt{32}} & 0 & 0 & 0 & 0 & 1 & -1 & 0
 \end{matrix} \begin{matrix}
 9 \\
 7 \\
 3 \\
 5 \\
 6 \\
 10 \\
 2 \\
 6
 \end{matrix}$$

Original image

the 1D Haar wavelet transform matrix

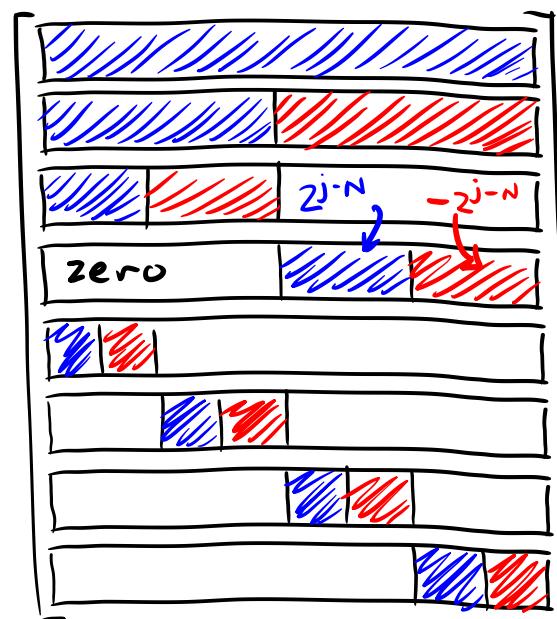
- Matrix contains $N-1$ scales
- Scale j represented by 2^j rows
 $\psi_0^j, \dots, \psi_{2^j-1}^j$



- Row ψ_j^i has $\frac{2^N}{2^j} = 2^{N-j}$ non-zero pixels
- They are pixels $x = i 2^{N-j}, \dots, (i+1)2^{N-j}-1$ with $|\psi_j^i(x)| = \frac{1}{2^{N-j}}$

Wavelet transformed image

$$\begin{bmatrix} I^0 \\ D^0 \\ D' \\ D^2 \end{bmatrix} = \begin{bmatrix} 6 \\ 0 \\ 2 \\ 2 \\ -1 \\ -1 \\ -2 \\ -2 \end{bmatrix}$$



$$\begin{bmatrix} 9 \\ 7 \\ 3 \\ 5 \\ 6 \\ 10 \\ 2 \\ 6 \end{bmatrix}$$

Original image